

Tail-Index Sure Independence Screening for High-Dimensional Heavy-Tailed Regression

Anonymous manuscript

Abstract

We propose a marginal screen for the covariates that determine a positive conditional tail index, designed for problems where the dimension p is large relative to the number of observations informative about the tail. The screen rests on one geometric fact: after projection onto a single coordinate, the conditional distribution is a mixture whose tail is governed by its heaviest component, so the projected tail index is the upper envelope of the tail-index surface along the corresponding fibre. Inactive coordinates produce flat envelopes at the global maximum γ^* , while an active coordinate whose fibres miss the global argmax set depresses the envelope; ranking coordinates by the average projected index therefore separates active variables from noise, and we characterize exactly which active coordinates a marginal screen can see, rather than assuming a generic signal-strength condition. We estimate the envelope with empirical-rank local Hill profiles. Ranks give every coordinate the same deterministic window geometry, which removes the covariate density conditions that fixed-dimension projected-tail theory requires and permits exact finite-sample union bounds over a diverging number of coordinates: the maximal score error over p coordinates is $O_{\mathbb{P}}(\sqrt{\log(pn)/(n\alpha h)})$ plus a bias term of order $o(\Delta_{\min})$, where α is the local tail fraction, h the rank bandwidth, $n\alpha h$ the effective local extreme sample size, and $\Delta_{\min}$ the weakest population gap. Sure screening and exact recovery follow when $\log(pn) = o(n\alpha h\Delta_{\min}^2)$, so a polynomially growing dimension is affordable whenever the local extreme count dominates the logarithm of the dimension. The algorithm costs $O(pn \log n)$ time and $O(n)$ working memory, under a second per screen at $(n, p) = (2000, 1000)$. A rank-aggregated version, taking the minimum rank over a block of nine neighbouring tunings, stabilizes the ranking against the choice of the smoothing parameters. Simulations across $p \in \{500, 1000, 2000\}$, including a model whose scale is driven by a correlated cluster of twenty inactive covariates, show when screening on the tail index differs from screening on a fixed high quantile: the cluster fills the leading positions of moderate-level quantile screens, whose contamination clears at extreme levels only at a severe variance cost, while the leading positions of the aggregated tail screen remain on the active set. An application to violent-crime rates in 1,993 U.S. communities, with $p = 99$ continuous predictors against an effective tail sample of about 200, illustrates the method in the regime it is designed for, and separates predictors of the limiting exponent from predictors of a high quantile on real data. The uniform projected-quantile conditions used by the theorems are derived from primitive assumptions on the tail-index function, the conditional tails and the fibre kernels, and verified for the simulation models.

1 Introduction

Let (X, Y) be a random pair with $Y > 0$ and $X \in \mathbb{R}^p$, and suppose that, conditionally on $X = x$,

$$\frac{\mathbb{P}(Y > ty \mid X = x)}{\mathbb{P}(Y > y \mid X = x)} \longrightarrow t^{-1/\gamma(x)}, \quad y \rightarrow \infty, \quad t > 0, \quad (1)$$

for a positive function γ , the conditional tail index. Estimating γ is the first step toward extreme conditional quantiles and tail probabilities [6, 12], and it suffers from a particularly severe curse of

dimensionality: local estimation must find observations near x and then keep only an upper fraction of their responses. A one-dimensional window of width h combined with a tail fraction α leaves an effective sample of order $n\alpha h$, not n [5, 10]. When p is large, even moderate sample sizes are effectively small.

This paper asks the screening question: among p candidate covariates, possibly with p growing with n , which ones does γ actually depend on? Sure independence screening [7] answers such questions by ranking coordinates through inexpensive marginal statistics and keeping a moderate number for subsequent joint analysis. Model-free variants exist for general dependence [13] and for fixed conditional quantiles [11], but the limiting exponent in (1) is a different target: a covariate can move the 0.95 quantile of Y without affecting γ , and conversely. Our simulations make this distinction concrete.

The screen we propose rests on a single geometric mechanism. Transform each continuous covariate to $U_j = F_j(X_j)$ and let $\xi_j(u)$ denote the tail index of Y given $U_j = u$ alone. Conditioning on one coordinate leaves the others free to range over the fibre $\{u\} \times [0, 1]^{p-1}$, so the projected distribution is a mixture, and a mixture of heavy tails is governed by its heaviest component. Under a full-support condition,

$$\xi_j(u) = \max_{v \in [0, 1]^{p-1}} \gamma(u, v) : \quad (2)$$

the projected index is the upper envelope of the tail-index surface along the fibre. If coordinate j is inactive, every fibre reaches the global maximum $\gamma^* = \max_x \gamma(x)$ and the envelope is flat at γ^* . If coordinate j is active and some of its fibres miss the global argmax set, the envelope dips below γ^* there. Averaging the envelope over a trimmed interval gives a score Ψ_j whose gap $\Delta_j = \gamma^* - \Psi_j$ is zero for every inactive coordinate and positive for every detectable active one. Ranking coordinates by increasing Ψ_j is the screen.

The same geometry states the method's limitation. An active coordinate is invisible when every one of its fibres still meets the global argmax set; this happens, for instance, for $\gamma(u_1, u_2) = \gamma_0 - (u_1 - u_2)^2$, whose argmax is the diagonal. Detectability is a property of the argmax projections, and we characterize it exactly rather than assuming a generic signal-strength condition.

Estimation follows the envelope. For each coordinate we slide a window over the empirical marginal ranks, apply a Hill statistic to the upper α -fraction of the responses in the window, and average the resulting profile. Empirical ranks give every coordinate the same deterministic window geometry and remove the marginal scales. The price is a rank-displacement term, which we control simultaneously over p coordinates by a Dvoretzky–Kiefer–Wolfowitz bound; it is negligible when $\log(pn) = o(nh^2)$. The main result bounds the maximal score error by

$$O_{\mathbb{P}} \left(\sqrt{\frac{\log(pn)}{n\alpha h}} \right) + o(\Delta_{\min}),$$

where $\Delta_{\min}$ is the weakest population gap and the second term collects the tail and spatial biases, and yields sure screening when $\log(pn) = o(n\alpha h \Delta_{\min}^2)$: high dimension is affordable exactly when the local extreme count dominates the logarithm of the dimension, so a polynomially growing p is admissible at a fixed gap. A blockwise implementation costs $O(pn \log n)$ time and $O(n)$ working memory beyond the scores.

Our construction adapts the projected-tail framework of Gardes and Podgorny [9], who estimate a low-dimensional central tail-index subspace by minimizing an average of projected local Hill estimates over projection matrices, in fixed dimension. We keep their projection semantics and local Hill estimator, and change what is needed for screening: the projections are the p coordinates themselves, p may diverge, windows are built from empirical ranks with deterministic local counts,

and uniformity over a compact matrix class is replaced by union bounds over coordinates and window states. Their density assumptions on the covariate are not needed here; the empirical-rank construction replaces them. Among screening proposals, Yoshida and Umezu [19] screen on covariate-dependent extreme-value indices through kernel conditional Pickands contrasts in a large-bandwidth regime, and parametric tail-index regression [18], extended to high dimensions with regularization by Sasaki et al. [17], exploits a parametric link when one is available. Tail dimension reduction methods [1, 2, 4, 8] target related but distinct conditional-tail objects.

The theoretical scope deserves a clear statement. The primitive model (1) identifies the projected index (2), while the uniform regularity needed for growing-dimensional concentration concerns the projected quantiles themselves, because projection creates logarithmic slow variation even from exact Pareto conditionals. We state the required conditions where they act, on the projected quantiles, and then derive them in Appendix C from primitive assumptions on γ , on the conditional tails and on the fibre kernels, after weakening the spatial condition to the range of quantiles the proofs actually use. The same appendix verifies those assumptions for the simulation models.

Section 2 develops the population theory. Section 3 defines the estimator and states the screening results. Section 4 presents the simulation models, the tuning study together with a rank-aggregated version of the screen built on a block of nine neighbouring tunings, and a comparison with competing screens at $p \in \{500, 1000, 2000\}$ that covers, in the same tables, the sensitivity to the dimension and to the quantile level of fixed-quantile screening. Section 5 applies the aggregated screen to violent-crime rates in U.S. communities. Section 6 discusses scope and practice. Proofs are in Appendix B, supported by the auxiliary results of Appendix A.

2 Population theory

2.1 Model and assumptions

All quantities below may depend on n : the dimension $p = p_n$, the tail-index function γ , the active set A , its size $s = |A|$, and later the bandwidth h , the tail fraction α and the retained size d . The index n is suppressed whenever no ambiguity arises.

Notation. A positive measurable function f is slowly varying at infinity if $f(ty)/f(y) \rightarrow 1$ as $y \rightarrow \infty$ for every $t > 0$, and a survival function $\bar{F}$ is regularly varying with index $-1/\xi < 0$ if $\bar{F}(ty)/\bar{F}(y) \rightarrow t^{-1/\xi}$ for every $t > 0$. Generalized inverses are $F^{-1}(t) = \inf\{y : F(y) \geq t\}$. For $u \in [0, 1]$ and $v \in [0, 1]^{p-1}$, $\iota_j(u, v) \in [0, 1]^p$ denotes the vector with u in coordinate j and v in the remaining coordinates, in order; we abbreviate $\gamma(u, v) = \gamma\{\iota_j(u, v)\}$ when j is clear from context. Norms on $[0, 1]^A$ are Euclidean. All asymptotic statements concern a triangular array indexed by n , and $Z_n = O_{\mathbb{P}}(a_n)$ means $\lim_{C \rightarrow \infty} \limsup_n \mathbb{P}(|Z_n| > Ca_n) = 0$.

The observed data are i.i.d. pairs (X_i, Y_i) with $Y_i > 0$ and $X_i \in \mathbb{R}^p$ with continuous marginal distribution functions $F_1, \dots, F_p$. The theory is formulated on the latent probability-integral transforms $U_{ij} = F_j(X_{ij})$, so that $U_i \in [0, 1]^p$ has uniform one-dimensional marginals and unrestricted cross-coordinate dependence; the transforms are never needed by the algorithm, because the ranks of $X_{1j}, \dots, X_{nj}$ coincide with the ranks of $U_{1j}, \dots, U_{nj}$ almost surely. With slight abuse of notation we write γ for the tail-index function on the cube, $\gamma(u) = \gamma_X\{F_1^{-1}(u_1), \dots, F_p^{-1}(u_p)\}$, where γ_X is the function appearing in (1).

(C1) *Conditional regular variation.* For every u , the conditional survival function factorizes as

$$\bar{F}(y \mid u) = \mathbb{P}(Y > y \mid U = u) = y^{-1/\gamma(u)} L(y, u),$$

where $L(\cdot, u)$ is slowly varying at infinity, $\gamma : [0, 1]^p \rightarrow (0, \infty)$ is continuous, and $0 < \underline{\gamma} \leq \gamma(u) \leq \Gamma < \infty$ uniformly in n and u .

(C2) *Uniform tail regularity.* There is a function c with $0 < c_- \leq c(u) \leq c_+ < \infty$ such that $\sup_u |L(y, u)/c(u) - 1| \rightarrow 0$ as $y \rightarrow \infty$.

(C1) states that each conditional law is heavy-tailed with index $\gamma(u)$, as in (1); (C2) strengthens the slow variation of L to uniform convergence toward a bounded positive limit, the form needed to control mixtures of the conditional laws.

The active set A collects the coordinates on which γ genuinely depends: $j \in A$ if and only if there exist u, v differing only in coordinate j with $\gamma(u) \neq \gamma(v)$. Iterating over the complementary coordinates shows that $\gamma(u) = g(u_A)$ for some function g , and that A is the smallest coordinate set with this property. Membership in A concerns the limiting exponent only, so a coordinate may shift the body of the conditional distribution, or its second-order behaviour, without being active. Conditions (C1)–(C2) parallel conditions (C.1)–(C.2) of Gardes and Podgorny [9], except that we do not assume the law of $\gamma(U)$ to be absolutely continuous; the projection argument below handles atoms, which arise as soon as one coordinate is inactive.

Fix a trimming parameter $0 < \varepsilon < 1/2$ and write $I_\varepsilon = [\varepsilon, 1 - \varepsilon]$. For each coordinate j , fix a pointwise conditional distribution $K_j(u, dv)$ of U_{-j} given $U_j = u$, and define the conditional law of Y given $U_j = u$ to be compatible with it through the disintegration identity

$$\mathbb{P}(Y > y \mid U_j = u) = \int \mathbb{P}\{Y > y \mid U = (u, v)\} K_j(u, dv);$$

without such pointwise versions the statements below hold for almost every u . Define

$$\xi_j(u) = \text{ess sup}_{v \sim K_j(u, \cdot)} \gamma(u, v),$$

the essential supremum of $v \mapsto \gamma(u, v)$ under $K_j(u, \cdot)$; Proposition 2.1 shows that the conditional law of Y given $U_j = u$ is regularly varying with index $-1/\xi_j(u)$, so $\xi_j(u)$ is its tail index. One further condition is used only where stated:

(S) *Full fibre support.* For every j and $u \in I_\varepsilon$, the support of $K_j(u, \cdot)$ is all of $[0, 1]^{p-1}$. A continuous joint density, positive on the interior cube, is sufficient.

2.2 Projection turns the surface into an envelope

Conditioning on $U_j = u$ mixes the conditional laws of Y over the fibre $\{u\} \times [0, 1]^{p-1}$. Heuristically, the slowest-decaying component of a mixture of heavy tails dictates the tail of the mixture, so the projected index should equal the largest tail index available on the fibre, and a localized region of heavy tails should survive projection intact. The following proposition makes this precise.

For a sub- σ -field $\mathcal{G}$, the conditional essential supremum $\text{ess sup}\{\gamma(U) \mid \mathcal{G}\}$ is the smallest $\mathcal{G}$ -measurable random variable dominating $\gamma(U)$ almost surely.

Proposition 2.1 (Projection). *Assume (C1)–(C2) and let $\mathcal{G} \subseteq \sigma(U)$. Then $\mathbb{P}(Y > y \mid \mathcal{G})$ is almost surely regularly varying with index $-1/\xi$, where $\xi = \text{ess sup}\{\gamma(U) \mid \mathcal{G}\}$. In particular $\xi_j(u)$ is the conditional essential supremum of $\gamma(U)$ given $U_j = u$. Under (S),*

$$\xi_j(u) = \max_{v \in [0, 1]^{p-1}} \gamma(u, v), \quad u \in I_\varepsilon. \quad (3)$$

The essential-supremum identity requires neither full support nor absolute continuity; (S) enters only to replace the essential supremum by the maximum over the whole fibre. If dependence removes part of a fibre, the maximum runs over the conditional support actually available at u , and the population target changes accordingly.




Figure 1: Projected envelopes. Left: a monotone active coordinate depresses the envelope below the inactive baseline γ^* . Right: in the diagonal example every fibre meets the argmax set, so the envelope of an active coordinate stays flat.

2.3 Score, gap, and detectability

Write $\gamma^* = \max_u \gamma(u)$ for the global maximum and

$$\Psi_j = \frac{1}{1-2\varepsilon} \int_{I_\varepsilon} \xi_j(u) du, \quad \Delta_j = \gamma^* - \Psi_j, \quad \Delta_{\min} = \min_{j \in A} \Delta_j. \quad (4)$$

The screen ranks coordinates by increasing Ψ_j and keeps the d smallest. Small scores indicate activity: all inactive coordinates share the flat baseline γ^* , while fixing a detectable active coordinate lowers the attainable envelope. The algorithm uses only the ordering of the Ψ_j and never estimates γ^* .

Let $\mathcal{M} = \arg \max_u \gamma(u)$, a nonempty compact set, and let $\pi_j(\mathcal{M})$ be its projection onto coordinate j .

Proposition 2.2 (Detectability). *Assume (C1)–(C2) and (S).*

1. If $j \notin A$, then $\xi_j \equiv \gamma^*$ on I_ε and $\Delta_j = 0$.
2. For every j and $u \in I_\varepsilon$: $\xi_j(u) = \gamma^*$ if and only if $u \in \pi_j(\mathcal{M})$.
3. $\Delta_j > 0$ if and only if $I_\varepsilon \setminus \pi_j(\mathcal{M})$ has positive Lebesgue measure.

If the restriction of γ to the active coordinates has a finite set of maximizers, every active coordinate is detectable.

Claim 3 separates two failure modes that finite-sample experiments tend to blur. When $\pi_j(\mathcal{M})$ covers I_ε , no sample size makes coordinate j identifiable from this score; the diagonal example $\gamma(u_1, u_2) = \gamma_0 - (u_1 - u_2)^2$, with $\mathcal{M} = \{(t, t) : 0 \leq t \leq 1\}$, is of this kind, and both of its active coordinates are invisible. When $\pi_j(\mathcal{M})$ misses a set of positive measure, detection is a question of estimation accuracy, quantified by Section 3. Figure 1 contrasts the two situations.

Proposition 2.3 (Quantitative margin). *Assume (C1)–(C2), (S), and that for some $u_A^* \in [0, 1]^A$ and $\nu > 0$,*

$$\gamma(u) \leq \gamma^* - c_1 \|u_A - u_A^*\|^\nu \quad \text{for all } u,$$

for some constant $c_1 > 0$. Then, for $j \in A$,

$$\gamma^* - \xi_j(u) \geq c_1 |u - u_j^*|^\nu, \quad \Delta_j \geq c_1 \mathbb{E}(|W - u_j^*|^\nu), \quad W \sim \text{Unif}(I_\varepsilon).$$

For $\nu = 2$ the right-hand side is at least $c_1\{(1 - 2\varepsilon)^2/12 + (1/2 - u_j^*)^2\}$.

Two consequences of (3) help with examples. If $\gamma(u) = G(\sum_{\ell \in A} g_\ell(u_\ell))$ with G strictly increasing, then $\xi_j(u) = G(g_j(u) + \sum_{\ell \in A, \ell \neq j} \max g_\ell)$, and $\Delta_j > 0$ whenever g_j lies below its maximum on a subset of I_ε of positive measure. If $\partial_j \gamma \geq c_j > 0$ throughout the cube, arbitrary interactions are allowed and $\Delta_j \geq c_j/2$.

For the benchmark $\gamma(u) = \frac{1}{2} \exp(-u_1 - u_2 - u_3 - u_4)$ used in Section 4, $\xi_j(u) = e^{-u}/2$ for $j \leq 4$ and $\xi_j \equiv 1/2$ otherwise, so $\Delta_j \approx 0.186$ at $\varepsilon = 0.05$. The gap is comfortable, yet estimating $\xi_j(u)$ near its maximum requires sample points whose remaining three active coordinates are simultaneously small; the simulations keep population detectability and finite-sample accessibility clearly separated.

3 Estimation and screening theory

3.1 Empirical-rank local Hill scores

The estimator follows the envelope construction coordinate by coordinate. Let R_{ij} be the rank of X_{ij} among $X_{1j}, \dots, X_{nj}$ and $\widehat{U}_{ij} = R_{ij}/(n+1)$. For $u \in I_\varepsilon$, form the rank window

$$\mathcal{W}_j(u) = \{i : |\widehat{U}_{ij} - u| \leq h\}, \quad M_j(u) = |\mathcal{W}_j(u)|, \quad k_j(u) = \lfloor \alpha M_j(u) \rfloor,$$

and, writing $Y_{j,u,(1)} \leq \dots \leq Y_{j,u,(M_j(u))}$ for the local order statistics, the local Hill estimate

$$\widehat{\xi}_j(u) = \frac{1}{k_j(u)} \sum_{r=1}^{k_j(u)} \log \frac{Y_{j,u,(M_j(u)-r+1)}}{Y_{j,u,(M_j(u)-k_j(u))}}, \quad (5)$$

defined when $\alpha M_j(u) > 1$. This is the local-fraction Hill statistic of Gardes and Podgorny [9, Definition 5] evaluated on one-coordinate rank windows. The score averages the profile over the deterministic grid $G = \{r/(n+1) : r = 1, \dots, n\} \cap I_\varepsilon$,

$$\widehat{\Psi}_j = \frac{1}{|G_j^{\text{val}}|} \sum_{u \in G_j^{\text{val}}} \widehat{\xi}_j(u), \quad G_j^{\text{val}} = \{u \in G : \alpha M_j(u) > 1\}, \quad (6)$$

Because the window counts $M_j(u)$ depend on the ranks only through the grid position, G_j^{val} is the same set for every j , and under the conditions of Theorem 3.1 every grid point is usable, so $G_j^{\text{val}} = G$; the convention $\widehat{\Psi}_j = +\infty$ when no grid point is usable is an implementation safeguard that is provably never triggered when $K_n \geq 2$. The screened set $\widehat{A}_d$ collects the d smallest scores, ties broken by a seeded auxiliary randomization.

The two localization parameters play different roles: h fixes how many observations resemble the target rank position, roughly $2nh$ of them, and α fixes how many of their largest responses enter the Hill statistic, roughly $2\alpha h$. Rank windows make these counts deterministic up to boundary effects, which is what permits exact finite-sample union bounds over a diverging number of coordinates.

3.2 Assumptions

Ranks displace the unobserved transforms U_{ij} , windows displace the conditioning point, and the Hill statistic sees only the upper α -fraction of a projected distribution. The assumptions bound these three effects uniformly over the coordinates. Throughout this section every active coordinate is assumed detectable, $\Delta_{\min} > 0$ for every n , so that the conditions below are well posed. Let $Q_j(\cdot, u)$ denote the generalized upper conditional quantile of Y given $U_j = u$, $Q_j(s, u) = F_{Y|j}^{-1}(1 - s^{-1} | u)$ with $F^{-1}(t) = \inf\{y : F(y) \geq t\}$, and let $\tau_n = (pn)^{-2}$.

(E1) *Quantile representation and regularity of the projected index.* For every $j \leq p$ and every $u \in [0, 1]$ the conditional distribution function $F_{Y|j}(\cdot | u)$ is continuous, so that $V_{ij} = 1 - F_{Y|j}(Y_i | U_{ij})$ is conditionally standard uniform and $Y_i = Q_j(1/V_{ij}, U_{ij})$ almost surely. Moreover the projected indices are equicontinuous on I_ε , uniformly over coordinates and over n ,

$$\omega_\xi(r) = \max_{j \leq p} \sup_{u, v \in I_\varepsilon, |u-v| \leq r} |\xi_j(u) - \xi_j(v)| \rightarrow 0 \quad (r \rightarrow 0),$$

and satisfy the grid-scale regularity condition $\omega_\xi(1/n) = o(\Delta_{\min})$.

(E2) *Uniform tail regularity.* Writing $Q_j(s, u) = s^{\xi_j(u)} \ell_j(s, u)$ with $\ell_j(\cdot, u)$ slowly varying,

$$\max_{j \leq p} \sup_{u \in I_\varepsilon} \sup_{s \geq (3\alpha)^{-1}, t > 1} \frac{|\log\{\ell_j(ts, u)/\ell_j(s, u)\}|}{\log t} = o(\Delta_{\min}).$$

(E3) *Uniform local quantile regularity.* There exists a constant $C_0 > 1$ such that

$$\max_{j \leq p} \sup_{u \in I_\varepsilon} \sup_{v \in [0, 1], |u-v| \leq C_0 h} \sup_{w \in [\tau_n, 1-\tau_n]} \left| \log \frac{Q_j(w^{-1}, v)}{Q_j(w^{-1}, u)} \right| = o(\Delta_{\min}).$$

The theorems use only this one C_0 ; under (7) any $C_0 > 1$ suffices.

Scale nuisances enter through (E3) and not through (E2), and the condition accommodates them. Writing $\log Q_j(s, u) = \xi_j(u) \log s + \log \ell_j(s, u)$, the displacement over $|v - u| \leq C_0 h$ splits into an index part, bounded by $2L_\xi C_0 h \log(pn)$ when ξ_j is Lipschitz with constant L_ξ , and a scale part which is $O(L_c C_0 h)$ whenever the conditional log-scale profile is Lipschitz in u with constant L_c uniform in n . A covariate that shifts the conditional scale of Y without touching the index is therefore harmless in the limit, at any bounded amplitude: both parts are $o(\Delta_{\min})$ as soon as $h \log(pn) = o(\Delta_{\min})$. Model B1 of Section 4 is built to make the finite-sample version of this statement measurable.

(E1) asks for continuity only; strict monotonicity of the conditional distribution function is not needed, because flat pieces are visited with probability zero. (E2) is the analogue of condition (H.4) in Gardes and Podgorny [9], normalized by $\log t$ and required only beyond the working tail level $(3\alpha)^{-1}$; it accommodates slowly varying factors such as $\ell(s) \asymp (\log s)^{-2}$, for which the left side is of order $1/\log(1/\alpha) \asymp 1/\log n$ when $\alpha = n^{-a}$, so gaps $\Delta_{\min} \gg 1/\log n$ are covered and smaller gaps are not. (E3) adapts their condition (H.5): the truncation $\tau_n = (pn)^{-2}$ makes the probability that any of the pn conditional PIT variables falls below τ_n vanish, which is the high-dimensional replacement for their fixed-dimension truncation. Their density conditions (H.1)–(H.2) have no counterpart here: empirical-rank windows have deterministic counts, so no lower bound on a covariate density is required.

These conditions are stated on the projected quantiles rather than on the primitive model, and deliberately so: projection generates slow variation of its own, which a condition imposed upstream

would have to anticipate. If $W \sim U(0, 1)$ and $\mathbb{P}(Y > y \mid W) = y^{-1/(1+W)}$ exactly, then $\mathbb{P}(Y > y) \sim 4y^{-1/2}/\log y$, so even exact Pareto conditionals produce a logarithmic factor after projection, and the left side of (E2) is then of exact order $1/\log(1/\alpha)$. Stating (E2) downstream measures that factor where it acts, and the resulting requirement is mild: every gap with $\Delta_{\min} \log(1/\alpha) \rightarrow \infty$ is covered, which for $\alpha = n^{-a}$ admits all gaps $\Delta_{\min} \gg 1/\log n$, and in particular every fixed gap.

Appendix C derives these conditions from primitive assumptions. Two steps make this possible. First, the range $w \in [\tau_n, 1 - \tau_n]$ in (E3) is wider than the proofs need: only the $k + 1$ largest responses of a window enter a Hill statistic, and Lemma A.12 shows that on the working event these are generated by $w \leq 3\alpha$. Replacing (E3) by the tail-only condition

$$\omega_n^\uparrow(C_0 h, \tau_n, \alpha) = \max_{j \leq p} \sup_{u \in I_\varepsilon} \sup_{|v-u| \leq C_0 h} \sup_{w \in [\tau_n, 3\alpha]} \left| \log \frac{Q_j(w^{-1}, v)}{Q_j(w^{-1}, u)} \right| = o(\Delta_{\min}) \quad (\text{E3}^\uparrow)$$

leaves Theorems 3.1 and 3.2 unchanged, and removes any requirement on the body of the conditional distribution. Second, Theorem C.3 derives (E1), (E2) and (E3[†]) from four primitive conditions — Lipschitz continuity and atomlessness, a differentiable uniform tail approximation, uniform doubling of the fibre-deficit mass for $D_{j,u} = 1/\gamma - 1/\xi_j(u)$, and a horizontal coupling of the primitive tails only — with

$$\omega_\xi(1/n) = O(n^{-1}), \quad b_n(3\alpha) = O\{1/\log(1/\alpha)\}, \quad \omega_n^\uparrow = O\{h \log(pn)\}.$$

Corollary C.7 verifies the primitive conditions for the Gaussian-copula models of Section 4, whose corner margins are regularly varying with a slowly varying Gaussian factor, so that a pure two-sided polynomial margin is neither assumed nor needed.

3.3 Main results

Theorem 3.1 (Score concentration). *Assume (C1)–(C2), (S), (E1)–(E2), (E3[†]), $\alpha \leq 1/3$, $h \leq 1 - \varepsilon$, and*

$$\log(pn) = o(nh^2), \quad \log(pn) = o(n\alpha h). \quad (7)$$

Then there are deterministic sequences $a_n = O(\sqrt{\log(pn)/(n\alpha h)})$ and $\beta_n = o(\Delta_{\min})$ with

$$\mathbb{P}\left(\max_{j \leq p} |\hat{\Psi}_j - \Psi_j| \leq a_n + \beta_n\right) \rightarrow 1;$$

in shorthand, $\max_{j \leq p} |\hat{\Psi}_j - \Psi_j| = O_{\mathbb{P}}(\sqrt{\log(pn)/(n\alpha h)}) + o(\Delta_{\min})$.

The stochastic term is the exponential-spacings fluctuation of $O(n)$ local Hill statistics per coordinate, each built from at least $c n \alpha h$ spacings, controlled simultaneously over p coordinates; $\sqrt{\log(pn)/(n\alpha h)}$ is the resulting effective-tail-sample-size rate. The $o(\Delta_{\min})$ term collects the tail bias (E2), the spatial bias (E3[†]) evaluated at $C_0 h$ — the first condition in (7) makes the rank displacement smaller than $(C_0 - 1)h$, so windows never leave the range covered by (E3[†]) — and the quadrature error $\omega_\xi(1/n) + O(1/n)$ from (E1).

Theorem 3.2 (Sure screening). *Assume the conditions of Theorem 3.1 (including $\Delta_{\min} > 0$, i.e. every active coordinate detectable in the sense of Proposition 2.2) and*

$$\log(pn) = o(n\alpha h \Delta_{\min}^2). \quad (8)$$

Then, for every $d \geq s$,

$$\mathbb{P}(A \subseteq \hat{A}_d) \rightarrow 1,$$

and with probability tending to one every active score is smaller than every inactive score.

Corollary 3.3 (Exact recovery). *Under the conditions of Theorem 3.2, if $d = s$ then $\mathbb{P}(\hat{A}_s = A) \rightarrow 1$.*

The proofs are in Appendix B, supported by the auxiliary results of Appendix A, whose centrepiece is a finite-sample uniform profile bound (Proposition A.14). The roadmap: (i) a Dvoretzky–Kiefer–Wolfowitz event per coordinate controls all empirical-rank displacements simultaneously, contributing the $\log(pn) = o(nh^2)$ requirement; (ii) rank windows have deterministic local counts and pass through at most $4n + 1$ membership states per coordinate; (iii) conditionally on a covariate column, the Rényi representation turns each local Hill statistic into a mean of i.i.d. exponentials built from about $n\alpha h$ extremes, plus a slowly varying remainder bounded through (E2); (iv) a union bound over the $p(4n + 1)$ states yields the $\log(pn)$ factor; (v) uniform score accuracy below $\Delta_{\min}/2$ gives deterministic separation of active from inactive scores. Gardes and Podgorny [9, Theorem 3] establishes the corresponding uniform consistency in fixed dimension over a compact class of projection matrices; their covering-number argument and density assumptions are replaced here by the rank construction, and their result does not imply the growing- p bound.

Condition (8) quantifies the cost of dimension: screening p coordinates costs $\log p$ against the effective local extreme count $n\alpha h$, deflated by the squared gap. With $\alpha = n^{-a}$ and $h = n^{-b}/2$ and $b \leq a$, dimensions up to $\log p = o(n^{1-a-b}\Delta_{\min}^2)$ are admissible, a polynomial range of p for fixed gaps; for $b > a$ the rank-displacement condition $\log(pn) = o(nh^2) = o(n^{1-2b})$ binds first. These ranges treat (E3[†]) as given; a spatial modulus of order $h \log(1/\tau_n)$, the natural benchmark, adds the constraint $\log(pn) = o(n^b \Delta_{\min})$.

4 Simulation study

Each coordinate is scored in $O(n \log n)$ time: one sort by empirical rank, then a sliding window over the rank grid whose local Hill statistics are maintained by a Fenwick tree indexed by global response ranks. Covariates are generated and processed in blocks, so the total cost is $O(pn \log n)$ time and $O(n)$ working memory beyond the scores. At $(n, p) = (2000, 1000)$ one complete screen takes under a second on a laptop core. Scores are sorted increasingly.

4.1 Simulation models

Family A: tail-index structure. Models A1, A2 and A3 share the covariate design, the slowly varying factor, and the tail-index active set $A_\gamma = \{1, 2, 3, 4\}$; they differ only in the tail-index function. Covariates are latent Gaussian AR(1),

$$Z_{i1} \sim N(0, 1), \quad Z_{ij} = \rho Z_{i,j-1} + \sqrt{1 - \rho^2} \epsilon_{ij}, \quad j = 2, \dots, p,$$

transformed to $U_{ij} = \Phi(Z_{ij})$; the tail-index functions are evaluated at the population transforms U , while the estimator sees only Z and computes its own empirical ranks. Responses are generated by quantile inversion, $Y = V^{-\gamma(U)} \ell(1/V, \cdot)$ with $V \sim \text{Unif}(0, 1)$, the mechanism of Gardes and Podgorny [9], and all three models carry their slowly varying factor

$$\ell_1(s, u) = [1 + \exp\{v(u) - s\}]^{-1}, \quad v(u) = 2.5(u_5 + u_6 + u_7 + u_8),$$

which satisfies $\ell_1(s, u) \rightarrow 1$ as $s \rightarrow \infty$: its effect vanishes in the limit while perturbing estimation at finite tail levels. The tail-index functions are

$$\begin{aligned} \gamma_{A1}(u) &= 0.5 \exp(-\sum_{j=1}^4 u_j), \\ \gamma_{A2}(u) &= 0.5 \exp(-0.5 \sum_{j=1}^4 u_j), \\ \gamma_{A3}(u) &= 0.5 \exp(-0.80 \sum_{j=1}^4 u_j - 0.20 u_1 u_2 - 0.15 u_3 u_4). \end{aligned}$$

A1 is the monotone reference, with trimmed population gap 0.186 (Section 2); its tail-index function and its Gaussian AR(1) covariate design are those of Yoshida and Umezu [19], adopted deliberately so that the comparison of Section 4.3 runs on a benchmark chosen by a competing method rather than on one of our own. The response mechanism differs: they invert a Burr-type conditional survival function, whereas A1–A3 carry the slowly varying factor ℓ_1 above. A2 weakens the signal, with gap 0.107. A3 adds interactions among the active coordinates while keeping a reasonably strong signal, with gap 0.158. All three share the maximum $\gamma^* = 0.5$.

Family B: finite-quantile scale contamination. Model B1 has the same tail-index function as A1, $\gamma_{B1} = \gamma_{A1}$, hence the same active set $A_\gamma = \{1, 2, 3, 4\}$ and the same gap 0.186, and *replaces* the slowly varying factor: responses are generated with ℓ_2 below in place of ℓ_1 , which does not appear in B1. Let $F_i \sim N(0, 1)$ be independent of the active coordinates and of the base AR(1) innovations, and set

$$\log \ell_2(s, F) = -\frac{1}{2s} + 0.5 \{ \Phi(F) - 0.5 \},$$

so the persistent scale factor lies in $[e^{-0.25}, e^{0.25}]$, strictly positive and uniformly bounded. The factor F is not observed. Instead, a block of $q_{\text{scale}} = 20$ proxies occupies the coordinates $A_{\text{scale}} = \{5, \dots, 24\}$:

$$Z_{ij} = 0.7 F_i + \sqrt{1 - 0.7^2} \epsilon_{ij}, \quad j \in A_{\text{scale}},$$

with independent standard normal proxy noises, so each proxy is marginally standard normal, loads on F at 0.7, and the pairwise correlation within the block is 0.49. The B1 covariate vector is built from a base AR(1) vector of length p : coordinates 1–4 and 25– p are kept from the base vector, coordinates 5–24 are replaced by the proxies, and the AR(1) chain beyond coordinate 24 continues from the base coordinate 24, not from the overwritten proxy. The observed B1 covariance is therefore deliberately not AR(1) over all coordinates; the AR(1) background is intact outside the proxy block. Because B1 changes both the slowly varying factor and the covariate design relative to A1, A1–B1 differences measure the joint effect of the scale cluster and of the exchange of ℓ_1 for ℓ_2 , not the cluster alone.

In B1, variables 1–4 determine the limiting tail index; variables 5–24 are informative about the finite conditional scale through their common information about F , and do not change the limiting exponent. A screen built on a fixed high quantile can therefore legitimately rank the proxy cluster highly — those variables move its target — even though the cluster is inactive for tail-index screening. B1 measures the practical consequence of this target difference. Section 3 places the mechanism in the theory: a bounded scale nuisance enters (E3) through a term of order $L_c h$, which vanishes with the bandwidth, so the screen is asymptotically immune to it; B1 measures how much sample that immunity costs at a finite bandwidth.

Monte Carlo design. The common design is

$$n = 2000, \quad p \in \{500, 1000, 2000\}, \quad \rho = 0.25, \quad d = 20, \quad \varepsilon = 0.05,$$

with 1,000 replications per cell and per model in the method comparison (binomial standard errors at most 0.016) and 200 common-random-number replications per model in the tuning study (standard errors at most 0.036); $\text{Sure-}d$ denotes $\mathbb{P}(R_{\max} \leq d)$ with $R_{\max} = \max_{j \in A_\gamma} \text{rank}(j)$, and all methods within a replication are evaluated on the same dataset.

Relation to the asymptotic theory. The simulation models instantiate the theory rather than sit beside it. Corollary C.7 verifies the primitive conditions of Appendix C for A1–A3 and B1, and shows that (E1), (E2) and (E3[†]) hold along the tuning sequences $\alpha = n^{-a}$, $h = n^{-b}/2$ used throughout, with bias orders $O(n^{-1})$, $O(1/\log n)$ and $O(n^{-b} \log n)$. The effective local extreme count is $n\alpha h = n^{1-a-b}/2$ while the population gaps are constants, so conditions (7) and (8) reduce to $b < 1/2$ and $a + b < 1$: the selected $(a, b) = (0.30, 0.15)$ and all nine settings of the aggregation block satisfy them, and Theorems 3.1 and 3.2 apply. The same corollary covers the aggregated screen, by intersecting the nine separation events.

What the finite design does not deliver is a numerically tight bound: at $n = 2000$ and $p = 1000$, $\sqrt{\log(pn)/(n\alpha h)} \approx 0.67$ against $\Delta_{\min} = 0.186$, and the lower endpoint entering (E2) is only $\log\{(3\alpha)^{-1}\} \approx 1.18$, so the asymptotic bounds are not yet informative at this sample size, as is usual for rate results at moderate n (Remark C.8). This is why the recovery rates below are reported as direct finite-sample evidence, and it is the practically relevant question: the theory says the screen eventually separates active from inactive coordinates, and the simulations measure how much sample the separation actually needs.

4.2 Tuning and the aggregated screen

The two smoothing parameters are parameterized as

$$\alpha = n^{-a}, \quad h = n^{-b}/2,$$

following the convention of Gardes and Podgorny [9]. Sure-20 is evaluated for each model over the crossed grid $a \in \{.25, .30, .35, .40, .45, .50\}$, $b \in \{0, .05, .10, .15, .20, .30, .40\}$ at $n = 2000$, $p = 1000$, $\rho = 0.25$, on 200 common-random-number replications per model: the 200 datasets are generated once per model and every cell of the grid is evaluated on the same datasets, so that differences between cells are not inflated by simulation noise. Figure 2 shows the four heatmaps.

The column $a = 0.25$ collapses for the three A-family models (Sure-20 at most .125, .080 and .225): at $\alpha = n^{-0.25} \approx 0.15$ the local Hill statistics reach into quantile regions where ℓ_1 , whose argument $v(u) - s$ is dominated by the U_5 – U_8 combination at moderate s , deforms the projected distributions. B1, whose bounded scale factor produces no such deformation, degrades only mildly there (up to .580). Performance also falls toward large a and large b together, where the local extreme count collapses, and the degenerate bandwidth $b = 0$ is never competitive. In between, the block $a \in \{0.30, 0.35\}$, $b \in \{0.10, 0.15, 0.20\}$ is the best-behaved region of the grid, but it is not a flat plateau. The four-model mean of Sure-20 is maximized at

$$(a^*, b^*) = (0.30, 0.15),$$

at 0.780; the five other cells of the block lie between 0.718 and 0.773, and the per-model values move by more than their paired standard errors across the block — A2, the weakest signal, falls from .555 at the argmax to .360 at (0.35, 0.20). The residual sensitivity within an otherwise good region is precisely what the aggregated screen below is designed to absorb. We keep the argmax as the single selected setting, with per-model Sure-20 .935, .555, .955 and .675; because it is selected on the same 200 replications shown in Figure 2, the comparison of Section 4.3 reruns everything on 1,000 independent replications. At $n = 2000$ the selected pair gives $\alpha \approx 0.10$, $h \approx 0.16$, and about $2n\alpha h = n^{1-a-b} \approx 65$ effective local extremes per interior window. No model is given its own tuning.

The aggregated screen. A single pair (a, b) makes the final ranking depend entirely on one tuning. The aggregated version of the screen uses the 3×3 block of settings

$$\mathcal{N}_9 = \{0.30, 0.35, 0.40\} \times \{0.10, 0.15, 0.20\},$$




Figure 2: Sure-20 over the (a, b) tuning grid for the four models at $n = 2000$, $p = 1000$, $\rho = 0.25$ (200 common-random-number replications per model; $\alpha = n^{-a}$, $h = n^{-b}/2$; standard errors at most 0.036). The dashed rectangle marks the aggregation block $\mathcal{N}_9$.

the neighbourhood of (a^*, b^*) inside the well-behaved region of the grid. For each setting $(a, b) \in \mathcal{N}_9$ and each coordinate j , compute the score $\hat{\Psi}_j(a, b)$ and its rank $r_j(a, b)$ among the p coordinates, ranks increasing with the score; aggregate by

$$r_j^{\text{agg}} = \min_{(a, b) \in \mathcal{N}_9} r_j(a, b),$$

and rank the coordinates by increasing r_j^{agg} ; coordinates sharing the same minimum are ordered by the median, then the mean, then the maximum of their nine ranks, and any remaining ties by a seeded auxiliary randomization. A coordinate is thus retained early whenever at least one reasonable tuning finds evidence for it. This direction of conservatism matches the purpose of a sure-screening first stage: false exclusion is final, while false retention is corrected by the second-stage multivariate analysis that the screen feeds. We make no optimality claim for the construction; it is the procedure used in the experiments and the application, and its effect is measured against the single-setting screen in Section 4.3. Its cost is nine score passes instead of one, leaving the $O(pn \log n)$ complexity unchanged up to the constant.

The two versions are reported side by side throughout: *Tail-index SIS, selected* (a^*, b^*) uses the central pair alone; *Tail-index SIS, aggregated* uses r_j^{agg} .

Table 1: Method comparison at $n = 2000$, $p = 500$, $\rho = 0.25$, on the same 1,000 replications per model. Sure- d is the probability that all four tail-index-active variables rank among the d smallest scores ($d = 4$ is exact recovery); $\mathbb{E}(R_{\max})$ and $\text{Med}(R_{\max})$ are the mean and median of the worst active rank $R_{\max} = \max_{j \in A_\gamma} \text{rank}(j)$. Standard errors of the probabilities are at most 0.016.

Model	Method	Sure-4	Sure-20	$\mathbb{E}(R_{\max})$	$\text{Med}(R_{\max})$
A1	Tail-index SIS, selected (a^*, b^*)	0.849	0.979	5.7	4
	Tail-index SIS, aggregated	0.884	0.990	4.6	4
	Yoshida–Umezū	0.202	0.467	97.4	29
	Quantile SIS ($\tau = .90$)	0.299	0.894	11.7	6
	Quantile SIS ($\tau = .95$)	0.891	0.998	4.2	4
	Quantile SIS ($\tau = .99$)	0.125	0.852	13.0	9
A2	Tail-index SIS, selected (a^*, b^*)	0.256	0.655	38.2	9
	Tail-index SIS, aggregated	0.196	0.797	20.3	8
	Yoshida–Umezū	0.028	0.178	153.9	124
	Quantile SIS ($\tau = .90$)	0.099	0.597	33.8	14
	Quantile SIS ($\tau = .95$)	0.431	0.930	8.2	5
	Quantile SIS ($\tau = .99$)	0.005	0.268	58.5	37
A3	Tail-index SIS, selected (a^*, b^*)	0.775	0.966	7.0	4
	Tail-index SIS, aggregated	0.790	0.986	5.1	4
	Yoshida–Umezū	0.181	0.509	79.8	20
	Quantile SIS ($\tau = .90$)	0.337	0.915	10.8	5
	Quantile SIS ($\tau = .95$)	0.863	0.996	4.4	4
	Quantile SIS ($\tau = .99$)	0.077	0.775	16.8	11
B1	Tail-index SIS, selected (a^*, b^*)	0.308	0.730	29.8	7
	Tail-index SIS, aggregated	0.448	0.907	9.8	5
	Yoshida–Umezū	0.007	0.088	132.0	108
	Quantile SIS ($\tau = .90$)	0.000	0.017	24.1	24
	Quantile SIS ($\tau = .95$)	0.030	0.514	18.0	20
	Quantile SIS ($\tau = .99$)	0.037	0.665	21.3	15

4.3 Comparison with competing screens

Six rules are compared on the same 1,000 simulated datasets per model and per dimension $p \in \{500, 1000, 2000\}$, at $n = 2000$, $\rho = 0.25$: the two versions of the tail-index screen; the conditional-Pickands screen of Yoshida and Umezū [19], implemented from their equations (2)–(4) with the Epanechnikov kernel and the tuning of their numerical work transplanted proportionally — their bandwidth $h = 1$ on standardized covariates is kept, and their extreme-value count is carried over as the same fraction of the sample size, $k = \lfloor 0.072n \rfloor$ — on 25 equally spaced evaluation points in $[\cdot 02, \cdot 98]$; and the quantile-adaptive SIS of He et al. [11] at the levels $\tau \in \{.90, .95, .99\}$, using cubic B-spline marginal quantile regression with three degrees of freedom, identical at every level. The competitors rank large utilities first; ours ranks small scores first. Tables 1–3 report, for each model and method, Sure-4 (exact recovery), Sure-20, and the mean and median of the worst active rank. Both competitors run at their authors’ settings, with no counterpart of the grid search of Section 4.2; in this respect the comparison favours the tail screen, and Section 4.2 shows how much its own performance moves across reasonable tunings.

Table 2: Method comparison at $n = 2000$, $p = 1000$, $\rho = 0.25$, on the same 1,000 replications per model. Sure- d is the probability that all four tail-index-active variables rank among the d smallest scores ($d = 4$ is exact recovery); $\mathbb{E}(R_{\max})$ and $\text{Med}(R_{\max})$ are the mean and median of the worst active rank $R_{\max} = \max_{j \in A_\gamma} \text{rank}(j)$. Standard errors of the probabilities are at most 0.016.

Model	Method	Sure-4	Sure-20	$\mathbb{E}(R_{\max})$	$\text{Med}(R_{\max})$
A1	Tail-index SIS, selected (a^*, b^*)	0.809	0.954	8.3	4
	Tail-index SIS, aggregated	0.823	0.985	5.1	4
	Yoshida–Umezū	0.179	0.398	188.1	58
	Quantile SIS ($\tau = .90$)	0.219	0.804	18.6	7
	Quantile SIS ($\tau = .95$)	0.830	0.997	4.6	4
	Quantile SIS ($\tau = .99$)	0.068	0.674	22.2	15
A2	Tail-index SIS, selected (a^*, b^*)	0.192	0.525	81.0	17
	Tail-index SIS, aggregated	0.151	0.677	39.5	10
	Yoshida–Umezū	0.021	0.138	320.1	256
	Quantile SIS ($\tau = .90$)	0.059	0.456	68.7	24
	Quantile SIS ($\tau = .95$)	0.303	0.873	14.8	6
	Quantile SIS ($\tau = .99$)	0.004	0.103	120.5	74
A3	Tail-index SIS, selected (a^*, b^*)	0.747	0.938	11.0	4
	Tail-index SIS, aggregated	0.706	0.980	6.7	4
	Yoshida–Umezū	0.176	0.445	156.8	34
	Quantile SIS ($\tau = .90$)	0.249	0.845	18.3	7
	Quantile SIS ($\tau = .95$)	0.794	0.992	4.7	4
	Quantile SIS ($\tau = .99$)	0.040	0.569	28.4	18
B1	Tail-index SIS, selected (a^*, b^*)	0.249	0.654	59.9	9
	Tail-index SIS, aggregated	0.368	0.849	17.0	5
	Yoshida–Umezū	0.004	0.059	265.5	206
	Quantile SIS ($\tau = .90$)	0.000	0.014	25.7	24
	Quantile SIS ($\tau = .95$)	0.035	0.489	18.7	21
	Quantile SIS ($\tau = .99$)	0.020	0.449	35.2	23

Sensitivity to the dimension. Every method degrades as p grows, at very different rates, and the mean–median contrast locates the degradation. For the selected-pair tail screen on A1, $\mathbb{E}(R_{\max})$ grows from 5.7 to 12.9 between $p = 500$ and $p = 2000$ while $\text{Med}(R_{\max})$ stays at 4: the typical replication is unaffected, and the loss is carried by a minority of replications in which an active coordinate falls into the null range, where its rank scales with p . Aggregation compresses that tail on every model ($\mathbb{E}(R_{\max})$ 4.6 $\rightarrow$ 6.2 on A1, 9.8 $\rightarrow$ 22.6 against 29.8 $\rightarrow$ 108.7 on B1, 5.1 $\rightarrow$ 8.1 against 7.0 $\rightarrow$ 22.0 on A3). The quantile screen at $\tau = .95$ is the least dimension-sensitive rule on the A family ($\mathbb{E}(R_{\max})$ 4.2 $\rightarrow$ 4.8 on A1). A2 is the most dimension-sensitive model for every method (aggregated Sure-20 .797 $\rightarrow$.561, quantile .95 .930 $\rightarrow$.782): a weak signal loses the most from added competition. Yoshida–Umezū degrades from a low base throughout ($\text{Med}(R_{\max})$ 29 $\rightarrow$ 168 on A1).

Sensitivity to the quantile level. On the A family the quantile screen is non-monotone in τ in every table, in opposite directions on either side of $\tau = .95$; on B1 its behaviour tracks its target, and Table 4 reports the resulting composition of its leading positions. At $\tau = .90$ its top

Table 3: Method comparison at $n = 2000$, $p = 2000$, $\rho = 0.25$, on the same 1,000 replications per model. Sure- d is the probability that all four tail-index-active variables rank among the d smallest scores ($d = 4$ is exact recovery); $\mathbb{E}(R_{\max})$ and $\text{Med}(R_{\max})$ are the mean and median of the worst active rank $R_{\max} = \max_{j \in A_\gamma} \text{rank}(j)$. Standard errors of the probabilities are at most 0.016.

Model	Method	Sure-4	Sure-20	$\mathbb{E}(R_{\max})$	$\text{Med}(R_{\max})$
A1	Tail-index SIS, selected (a^*, b^*)	0.751	0.941	12.9	4
	Tail-index SIS, aggregated	0.802	0.981	6.2	4
	Yoshida–Umezū	0.125	0.305	432.2	168
	Quantile SIS ($\tau = .90$)	0.161	0.665	36.6	11
	Quantile SIS ($\tau = .95$)	0.758	0.992	4.8	4
	Quantile SIS ($\tau = .99$)	0.025	0.406	40.2	25
A2	Tail-index SIS, selected (a^*, b^*)	0.149	0.463	144.8	27
	Tail-index SIS, aggregated	0.088	0.561	76.9	15
	Yoshida–Umezū	0.015	0.084	640.2	545
	Quantile SIS ($\tau = .90$)	0.039	0.328	122.2	42
	Quantile SIS ($\tau = .95$)	0.234	0.782	24.7	7
	Quantile SIS ($\tau = .99$)	0.000	0.022	245.7	154
A3	Tail-index SIS, selected (a^*, b^*)	0.650	0.892	22.0	4
	Tail-index SIS, aggregated	0.657	0.963	8.1	4
	Yoshida–Umezū	0.142	0.379	336.7	67
	Quantile SIS ($\tau = .90$)	0.171	0.723	36.4	9
	Quantile SIS ($\tau = .95$)	0.698	0.988	5.4	4
	Quantile SIS ($\tau = .99$)	0.023	0.328	56.2	34
B1	Tail-index SIS, selected (a^*, b^*)	0.234	0.599	108.7	11
	Tail-index SIS, aggregated	0.339	0.805	22.6	6
	Yoshida–Umezū	0.008	0.057	494.5	380
	Quantile SIS ($\tau = .90$)	0.000	0.012	25.1	24
	Quantile SIS ($\tau = .95$)	0.038	0.482	19.0	21
	Quantile SIS ($\tau = .99$)	0.008	0.254	61.7	39

four contain on average 3.8 of the twenty scale proxies and all twenty fill its top 24; exact recovery is zero to three decimals at every p . This is coherent behaviour for a .90-quantile target, which the proxies genuinely move, and it is fatal for tail-index recovery. At $\tau = .95$ the contamination remains substantial (2.4–2.5 proxies in the top four, 19.1–19.5 in the top 24, median worst active rank 20–21). At $\tau = .99$ the contamination largely clears (0.2–0.4 proxies in the top four) because the quantile target approaches the tail index, but the variance cost is severe everywhere: on the A family, Sure-4 falls from .234–.891 at $\tau = .95$ to at most .125 at $\tau = .99$, and on B1 the level .99 never reaches Sure-4 .04. Yoshida–Umezū, built on Pickands-type quantile differences at its own fixed levels, is dominated on all models and dimensions here, with 6.8–9.8 of the twenty B1 proxies in its top 24 (Table 4).

Effect of the aggregation. Aggregation improves the tail screen most where failures are worst. On B1 it lifts exact recovery from .308/.249/.234 to .448/.368/.339 across the three dimensions, divides $\mathbb{E}(R_{\max})$ by three to five, and brings $\text{Med}(R_{\max})$ to 5–6; on A1 and A3 it moves Sure-4 by a

Table 4: Model B1: average number of the twenty scale proxies among the top 4 and top 24 positions of each ranking, by dimension (same 1,000 replications as Tables 1–3). A perfectly specific screen scores 0; ranking the proxies first scores 4 and 20.

Method	$p = 500$		$p = 1000$		$p = 2000$	
	top 4	top 24	top 4	top 24	top 4	top 24
Tail-index SIS, selected (a^*, b^*)	0.32	2.73	0.31	2.30	0.27	1.72
Tail-index SIS, aggregated	0.19	2.71	0.15	2.15	0.12	1.57
Yoshida–Umezū	1.78	9.84	1.50	8.03	1.35	6.82
Quantile SIS ($\tau = .90$)	3.82	20.00	3.82	20.00	3.79	20.00
Quantile SIS ($\tau = .95$)	2.48	19.53	2.48	19.38	2.43	19.05
Quantile SIS ($\tau = .99$)	0.44	4.53	0.32	3.06	0.21	1.85

few points in either direction (+3.5, +1.4, +5.1 on A1; +1.5, −4.1, +0.7 on A3) while compressing the rank tail; Sure-20 improves in all twelve model–dimension cells. On the weak-signal A2 it trades exact recovery (.256 → .196 at $p = 500$, .149 → .088 at $p = 2000$) for substantially better inclusion (Sure-20 .463 → .561 at $p = 2000$; $\mathbb{E}(R_{\max})$ 144.8 → 76.9). The trade is the intended one: promotion by a single reasonable tuning occasionally admits a lucky null into the top four when the signal is weak, and reliably rescues active coordinates from single-setting failures further down the list. Aggregation stabilizes inclusion — the sure-screening objective — rather than exact selection.

The quantile screen at a moderate level remains the strongest detector when high quantiles and the limiting exponent share their covariates and the signal is weak (A2); the aggregated tail screen is competitive or better on A1 and A3, and on B1 it is the only rule whose leading positions stay on the active set — Sure-4 .448/.368/.339 and Sure-20 .907/.849/.805 across the three dimensions, against at most .038 and .665 for any fixed-quantile level, with 0.12–0.19 proxies in its top four (Table 4). It is degraded, not immune: the same rule reaches Sure-4 .884/.823/.802 on A1, so the scale cluster, jointly with the exchange of slowly varying factors, costs it about forty-five points of exact recovery. For a question about the limiting exponent itself, the two rankings are best read side by side, as in Section 5: variables high on one list and low on the other are where the two targets part ways.

5 Application: violent crime across U.S. communities

The analysis contains $n = 1,993$ communities and $p = 99$ candidate predictors. Under the tuning selected in Section 4.2, $\alpha = n^{-0.3} \approx 0.102$ and $h = n^{-0.15}/2 \approx 0.160$, so the sample contains $n\alpha \approx 204$ upper-tail observations in total and about $2nah \approx 65$ within a typical interior rank window. The dimension is therefore comparable to the number of observations available for local tail estimation. Unrestricted or nonparametric joint tail modelling with 99 covariates is poorly supported by roughly 200 upper-tail observations, and by fewer still within any local covariate neighbourhood; a marginal screen can rank the 99 candidates from exactly this information.

5.1 Data and protocol

The Communities and Crime data set [15], distributed by the UCI Machine Learning Repository in its unnormalized version, links socio-economic census attributes, law-enforcement survey attributes and FBI crime counts for 2,215 U.S. communities. The response is the violent-crime rate per 100,000 inhabitants. We keep the communities with an observed positive response ($n = 1,993$). Yoshida and

Umezū [19] analyse the same data set with their own screen, taking the robbery rate as response; we deliberately stay on the same data so that the two screens can be read side by side, and take the violent-crime rate, whose larger positive count leaves a usable upper tail after the response-blind filtering described below. Preprocessing of the 124 documented predictive attributes is response-blind: we retain numeric predictors with fewer than 5% missing values, which removes the 22 law-enforcement variables reported by only 343 communities, and with at least 100 distinct observed values, which removes 3 low-cardinality variables, keeping the application close to the continuous-marginal framework of the theory. This leaves $p = 99$ continuous predictors — demographic composition, income and poverty, education, employment, family structure, immigration, language, housing, and population density. The at most one missing value per retained predictor is median-imputed, and every method below works from the same imputed matrix.

The response is strongly right-skewed: median 374, 99th percentile 2,943, maximum 4,877. The left panel of Figure 3 shows the Hill curve of the response over $k \in [20, 600]$ upper order statistics; it is positive and moves slowly, between 0.283 and 0.374 for $k \in [100, 300]$, with value 0.328 at the working level $k = n\alpha = 204$. The middle panel shows the Pareto QQ plot of the 600 largest responses, approximately linear over the working range with least-squares slope 0.27. The gap between the QQ slope and the Hill value at the same level is mild evidence that regular variation is only approximate over the working range. The response exhibits a sufficiently heavy and stable upper tail to make a tail-index analysis plausible; neither diagnostic establishes regular variation beyond doubt, and no single threshold carries a structural interpretation.

The screening protocol is the one fixed in Section 4.2, applied unchanged: the aggregated tail-index screen computes the score at each of the nine settings $\mathcal{N}_9 = \{0.30, 0.35, 0.40\} \times \{0.10, 0.15, 0.20\}$, converts each pass into ranks, and orders the predictors by the minimum rank r_j^{agg} across the nine settings; the selected-pair screen at $(a^*, b^*) = (0.30, 0.15)$ is reported alongside. The two competing screens of Section 4.3 are run on the same imputed matrix: Yoshida–Umezū at its transplanted tuning and the quantile screen at $\tau \in \{.90, .95, .99\}$.

5.2 Results

Poverty and family structure dominate the aggregated ranking, and they dominate every competing ranking as well. The two-parent household variable `pctKids2Par` is first (per-setting ranks 1–4 across the nine settings), the poverty rate `pctPoverty` second (ranks 1–6) and `pct2Par` third (ranks 1–5); the two family-structure variables are also leading variables for the quantile screen at every level (`pctKids2Par` is first at all three τ) and for Yoshida–Umezū (ranks 3–7). The related `pctKids4w2Par` and `pct12to17w2Par` complete the consensus block at aggregated ranks 7–8. The high-quantile and tail-index rankings share these leading predictors.

The informative disagreements are the tail-specific candidates. The share of workers using public transportation (`pctUsePubTrans`, aggregated rank 5) and the farm-employment rate (`pctWfarm`, rank 6) rank far lower under every fixed-quantile method: ranks 49–73 for the quantile screen across its three levels and 57–60 for Yoshida–Umezū. Their low tail-index scores are consistent with marginal variation in upper-tail heaviness across their values, of a kind that fixed-quantile screens are not designed to detect; Section 4.3 produces the same dissociation under controlled conditions, where the target is known. The two variables differ in stability across the aggregation block: `pctUsePubTrans` ranks within the top 15 at every one of the nine settings, while `pctWfarm` moves between 2 and 47 (right panel of Figure 3). The never-married male rate (`pctMaleNevMar`, aggregated rank 4, per-setting ranks 1–9, quantile ranks 28–35) sits between the consensus block and the tail-specific pair. We report all three as candidate tail-specific predictors for a second-stage analysis.




Figure 3: Communities and Crime application ($n = 1,993$, $p = 99$, $\alpha = n^{-0.3}$, $h = n^{-0.15}/2$). Left: Hill estimates for the violent-crime rate; the dashed line marks the working level $k = n\alpha = 204$. Middle: Pareto QQ plot of the 600 largest responses; the dashed line is fitted over the working range $i \leq 204$. Right: for the ten leading predictors of the aggregated screen, the nine per-setting ranks (open circles) and the aggregated rank (filled circles); the dotted line is the default retained size $d = 20$.

Minimum-rank aggregation is deliberately conservative: a predictor is promoted whenever at least one tuning in the prespecified neighbourhood ranks it highly, which suits a first-stage sure screen, where false exclusion is more costly than retaining additional candidates. **pctHousW0plumb** (aggregated rank 9) and **pctLargHous** (rank 10) illustrate the resulting behaviour: their positions rest on one favourable setting each (per-setting ranks 6–55 and 6–74; selected-pair ranks 36 and 46), so their rankings are markedly more tuning-sensitive than those of the consensus block or of the tail-specific candidates, which hold their positions across the whole block. The conservative screen retains them intentionally; the per-setting ranks describe this sensitivity, and discrimination among the retained predictors is left to the subsequent multivariate analysis.

We do not interpret any ranking causally: the screen identifies covariates along which the conditional tail index varies marginally, and confounding among the 99 predictors is certain.

6 Discussion

The screen is a first stage. Its intended pipeline reduces p coordinates to a manageable d , after which a joint method — the central tail-index subspace of Gardes and Podgorny [9], a structured tail regression [17, 18], or multivariate local estimation — can resolve interactions and correlated neighbours among the retained variables. The two stages solve different problems: rapid elimination of evidently irrelevant coordinates scales to large p , while accurate joint tail inference does not.

The scope of a marginal screen for this target is delimited exactly rather than assumed. An active coordinate every one of whose fibres meets the global argmax set leaves no trace in the score, at any sample size; Proposition 2.2 characterizes this case exactly, in terms of the argmax projections, where a signal-strength condition would only have hidden it. Knowing the boundary is what allows a user to tell a screen that has failed from a target that a marginal screen cannot see, and it marks where iterative or grouped extensions, with their own identification arguments, would

be required. Two further features of the target are worth stating in the same spirit. Projected mixture bias is logarithmic, so (E2) covers every gap with $\Delta_{\min} \log(1/\alpha) \rightarrow \infty$ and in particular every fixed gap; whether gaps of the excluded order $1/\log n$ are detectable by any estimator is a question about the target, not about this screen. And the envelope near its maximum is estimated from observations whose remaining active coordinates are simultaneously near their maximizing values, so the sample cost grows with the number of active coordinates: the weak-signal model A2 in Section 4 measures it, and aggregation over neighbouring tunings recovers most of it.

Appendix C derives the estimation conditions from primitive assumptions that are verifiable on concrete models; what remains is to establish whether the logarithmic projected bias, which restricts the covered gaps to $\Delta_{\min} \log(1/\alpha) \rightarrow \infty$, is a property of the target or of this estimator. The score functional also invites variants: a minimum or a low quantile of the profile would emphasize depressions confined to short intervals, at the price of new uniformity arguments for empirical profile extrema.

In practice, tuning matters more than the theory prescribes. The tail fraction trades Hill bias against variance; the bandwidth additionally determines how often near-maximizing configurations of the other active coordinates enter a window. We recommend the workflow of Sections 4 and 5: run the screen on the block of nine neighbouring tunings and rank by the aggregated minimum rank. The aggregation is deliberately conservative — a predictor is retained when at least one reasonable tuning finds evidence for it — which is the right direction of error for a first-stage sure screen: false exclusion is final, while false retention is corrected by the second-stage multivariate analysis. The per-setting ranks describe tuning sensitivity and should accompany the ranking as a diagnostic, not act as a second screening threshold. Dependence deserves caution of its own: it can help, by concentrating active configurations, or hurt, by promoting inactive neighbours, and the direction is design-specific.

Projection turns a tail-index surface into an upper envelope. Variation of that envelope is a scalable marginal signal for tail-index activity, and its absence, characterized by the argmax projections, tells the analyst exactly what a marginal screen can and cannot see.

A Auxiliary results

This appendix collects the technical tools used by the proofs of Appendix B: elementary probabilistic lemmas (A.1), the projection lemmas behind Proposition 2.1 (A.2), the deterministic geometry of empirical-rank windows (A.3), the local Hill representation and the finite-sample uniform profile bound (A.4), and the passage from profiles to scores, ending with a deterministic separation lemma (A.5).

Throughout, the dimension is written $p = p_n$, and two auxiliary moduli quantify the content of (E2)–(E3):

$$b_n(a) = \max_{j \leq p} \sup_{u \in I_\varepsilon} \sup_{s \geq a^{-1}, t > 1} \frac{|\log\{\ell_j(ts, u)/\ell_j(s, u)\}|}{\log t}, \quad \omega_n(r, \tau) = \max_{j \leq p} \sup_{u \in I_\varepsilon} \sup_{\substack{v \in [0, 1] \\ |v - u| \leq r}} \sup_{w \in [\tau, 1 - \tau]} \left| \log \frac{Q_j(w^{-1}, v)}{Q_j(w^{-1}, u)} \right|,$$

so that (E2) reads $b_n(3\alpha) = o(\Delta_{\min})$ and (E3) reads $\omega_n(C_0 h, \tau_n) = o(\Delta_{\min})$ with $\tau_n = (pn)^{-2}$; $\omega_n(\cdot, \tau)$ is nondecreasing in its first argument, so (E3) bounds $\omega_n(r, \tau_n)$ for every $r \leq C_0 h$.

The proofs below never use the whole range $w \in [\tau, 1 - \tau]$: only the $k + 1$ largest responses of a window enter a Hill statistic, and Lemma A.12 shows that on the working event these correspond to

$w \leq 3\alpha$. We therefore also write, for $0 < \tau \leq 3\alpha < 1$,

$$\omega_n^\uparrow(r, \tau, \alpha) = \max_{j \leq p} \sup_{u \in I_\epsilon} \sup_{\substack{v \in [0,1] \\ |v-u| \leq r}} \sup_{w \in [\tau, 3\alpha]} \left| \log \frac{Q_j(w^{-1}, v)}{Q_j(w^{-1}, u)} \right| \leq \omega_n(r, \tau), \quad (\text{A.1})$$

and record the tail-only condition

$$\omega_n^\uparrow(C_0 h, \tau_n, \alpha) = o(\Delta_{\min}), \quad \tau_n = (pn)^{-2}, \quad (\text{E3}^\uparrow)$$

which is implied by (E3) and suffices everywhere below. Every result of this appendix and of Appendix B is stated with $\omega_n^\uparrow$; Appendix C gives primitive sufficient conditions for (E1), (E2) and (E3[†]). Put $\delta_n = (n+1)^{-1}$,

$$m_n^- = \lfloor \{h + \min(h, \epsilon)\}(n+1) \rfloor - 1, \quad K_n = \lfloor \alpha m_n^- \rfloor, \quad S_n = 4n + 1.$$

All conditional distributions are the pointwise compatible versions of Section 2, and constants are uniform over the triangular array.

A.1 Elementary probabilistic lemmas

The first lemma transfers pointwise perturbations to order statistics, as in Gardes and Podgorny [9, Appendix, Lemma 3].

Lemma A.1 (Sorting is sup-norm nonexpansive). *Let $a_1, \dots, a_m$ and $b_1, \dots, b_m$ satisfy $a_i \leq b_i$ for all i . Then $a_{(i)} \leq b_{(i)}$ for all i . Consequently $|a_i - b_i| \leq \epsilon$ for all i implies $|a_{(i)} - b_{(i)}| \leq \epsilon$ for all i .*

Proof. Fix i . At least $m - i + 1$ indices r satisfy $a_r \geq a_{(i)}$, and for each such r , $b_r \geq a_r \geq a_{(i)}$; hence $b_{(i)} \geq a_{(i)}$. For the second claim apply the first to $(a_i - \epsilon, b_i)$ and $(b_i - \epsilon, a_i)$. $\square$

Lemma A.2 (Chernoff bound for exponential means). *Let L_k be the mean of k i.i.d. standard exponentials. For $0 < t \leq 1$, $\mathbb{P}(|L_k - 1| > t) \leq 2e^{-kt^2/4}$.*

Proof. Markov's inequality applied to $e^{\lambda k L_k}$, optimized at $\lambda = t/(1+t)$, gives $\mathbb{P}(L_k > 1+t) \leq e^{-k\{t - \log(1+t)\}}$. The function $\varphi(t) = t - \log(1+t) - t^2/4$ has $\varphi(0) = 0$ and

$$\varphi'(t) = 1 - \frac{1}{1+t} - \frac{t}{2} = \frac{t(1-t)}{2(1+t)} \geq 0 \quad \text{on } [0, 1],$$

so $t - \log(1+t) \geq t^2/4$ on $(0, 1]$. For the lower tail, $\mathbb{P}(L_k < 1-t) \leq e^{-k\{-t - \log(1-t)\}}$ and $-t - \log(1-t) \geq t^2/2 \geq t^2/4$. $\square$

Lemma A.3 (Threshold order statistic of uniforms). *Let $V_{(1)} < \dots < V_{(m)}$ be the order statistics of m i.i.d. standard uniforms, $0 < \alpha \leq 1/3$ and $k = \lfloor \alpha m \rfloor \geq 1$. Then $\mathbb{P}\{V_{(k+1)} > 3\alpha\} \leq e^{-(k+1)/4}$.*

Proof. If $3\alpha \geq 1$ the probability is zero. Otherwise $\{V_{(k+1)} > 3\alpha\} = \{N \leq k\}$ with $N \sim \text{Bin}(m, 3\alpha)$ and mean $\mu = 3\alpha m \geq 3k$. The multiplicative Chernoff bound $\mathbb{P}\{N \leq (1-\delta)\mu\} \leq e^{-\delta^2\mu/2}$ [3, Section 2.2] with $\delta = 2/3$ gives $\mathbb{P}(N \leq \mu/3) \leq e^{-2\mu/9} \leq e^{-2k/3} \leq e^{-(k+1)/4}$, the last step because $8k \geq 3(k+1)$ for $k \geq 1$. $\square$

A.2 Projection lemmas

Let $\mathcal{G} \subseteq \sigma(U)$, $T = \gamma(U)$, and $\xi = \text{ess sup}(T \mid \mathcal{G})$, the smallest $\mathcal{G}$ -measurable random variable dominating T almost surely.

Lemma A.4 (Uniform remainder replacement). *Define $N(y) = \mathbb{E}\{c(U)y^{-1/T} \mid \mathcal{G}\}$. Under (C1)–(C2), almost surely,*

$$\left| \frac{\mathbb{P}(Y > y \mid \mathcal{G})}{N(y)} - 1 \right| \leq \rho(y) := \sup_u \left| \frac{L(y, u)}{c(u)} - 1 \right| \longrightarrow 0. \quad (\text{A.2})$$

Proof. By the tower property, $\mathbb{P}(Y > y \mid \mathcal{G}) = \mathbb{E}\{y^{-1/T} L(y, U) \mid \mathcal{G}\}$. Writing $L(y, U) = c(U)\{1 + r(y, U)\}$ with $|r(y, U)| \leq \rho(y)$,

$$|\mathbb{P}(Y > y \mid \mathcal{G}) - N(y)| = |\mathbb{E}\{c(U)y^{-1/T} r(y, U) \mid \mathcal{G}\}| \leq \rho(y)N(y),$$

and $\rho(y) \rightarrow 0$ is (C2). $\square$

The next lemma is the point at which the argument departs from the proof of Proposition 1 in Gardes and Podgorny [9]: their condition (C.1) assumes the law of $\gamma(X)$ absolutely continuous, which identifies $\{\gamma > t\}$ with $\{\gamma \geq t\}$ in their computation. The tilting argument below needs no such assumption and allows purely atomic conditional laws of T — the relevant case for inactive coordinates and constant-index null models.

Lemma A.5 (Tilted-measure concentration). *For $y > 1$ define the random probability measure*

$$K_y^*(A) = \frac{\mathbb{E}\{c(U)y^{-1/T} \mathbf{1}_{\{T \in A\}} \mid \mathcal{G}\}}{N(y)}, \quad A \in \mathcal{B}([\underline{\gamma}, \Gamma]).$$

On a common probability-one event, $K_y^ \Rightarrow \delta_\xi$ weakly as $y \rightarrow \infty$.*

Proof. Since $K_y^*\{T > \xi + \delta\} = 0$ for every $\delta > 0$ by definition of ξ , it suffices to prove, for fixed $\delta \in (0, \underline{\gamma})$,

$$K_y^*\{T \leq \xi - \delta\} \longrightarrow 0 \quad \text{almost surely.} \quad (\text{A.3})$$

Split by the $\mathcal{G}$ -measurable event $G_\delta = \{\xi \geq \underline{\gamma} + \delta\}$. On G_δ^c , $\xi - \delta < \underline{\gamma} \leq T$, so $\{T \leq \xi - \delta\}$ is empty and (A.3) is trivial. On G_δ , set $r = \xi - \delta$, $s = \xi - \delta/2$ and $\pi = \mathbb{P}(T > s \mid \mathcal{G})$. Then $\pi > 0$ almost surely on G_δ : otherwise $\xi' = \xi - (\delta/2)\mathbf{1}_{G_\delta \cap \{\pi=0\}}$ would be a $\mathcal{G}$ -measurable almost-sure upper bound for T , strictly smaller than ξ with positive probability, contradicting minimality of the conditional essential supremum. Bounding numerator and denominator separately, using only $c_- \leq c \leq c_+$ and monotonicity of $g \mapsto y^{-1/g}$,

$$K_y^*\{T \leq r\} \leq \frac{c_+ y^{-1/r}}{c_- \pi y^{-1/s}} = \frac{c_+}{c_- \pi} y^{-(1/r - 1/s)} \longrightarrow 0.$$

A rational sequence $\delta \downarrow 0$ and a countable intersection give a common probability-one event. $\square$

Lemma A.6 (Continuity of the envelope). *Under (C1) and (S), so that $\xi_j(u)$ equals the fibre maximum $\max_v \gamma(u, v)$, each ξ_j is continuous on $[0, 1]$, with modulus bounded by the modulus of continuity of γ .*

Proof. γ is uniformly continuous on the compact cube; if $|u - u'| \leq \delta$ then $|\gamma(u, v) - \gamma(u', v)| \leq \varpi(\delta)$ for every v , and taking maxima preserves the bound. $\square$

A.3 Empirical-rank geometry

The three lemmas of this subsection are purely deterministic or rank-based; no property of the response enters. They replace the covariate-density conditions (H.1)–(H.2) of Gardes and Podgorny [9] in the present empirical-rank setting.

Lemma A.7 (Rank displacement). *Let $\widehat{U}_{ij} = R_{ij}/(n+1)$ and*

$$D_n(\eta) = \left\{ \max_{j \leq p} \max_{i \leq n} |U_{ij} - \widehat{U}_{ij}| \leq \eta + \delta_n \right\}.$$

Then $\mathbb{P}\{D_n(\eta)^c\} \leq 2pe^{-2n\eta^2}$.

Proof. Fix j and let $F_{n,j}$ be the empirical distribution function of the U_{ij} , $i \leq n$. By the Dvoretzky–Kiefer–Wolfowitz inequality with Massart’s constant [14], $\sup_t |F_{n,j}(t) - t| \leq \eta$ outside probability $2e^{-2n\eta^2}$. On that event, continuity gives $F_{n,j}(U_{(r)j}) = r/n$, hence $|U_{(r)j} - r/n| \leq \eta$ and

$$|U_{(r)j} - r/(n+1)| \leq \eta + \frac{r}{n(n+1)} \leq \eta + \delta_n$$

for every rank r simultaneously. A union bound over j finishes; no union over ranks is needed. $\square$

Lemma A.8 (Deterministic window counts). *For every j , every realization and every $u \in I_\varepsilon$, $M_j(u) \geq m_n^-$ and $k_j(u) \geq K_n$; moreover $M_j(u)$ is a nonrandom function of u .*

Proof. The ranks form a permutation of $1, \dots, n$, so $M_j(u) = \#\{r \in \{1, \dots, n\} : |r/(n+1) - u| \leq h\}$, a deterministic function of u : the number of integers in $[x(u), y(u)]$ with $x(u) = \max\{(u-h)(n+1), 1\}$ and $y(u) = \min\{(u+h)(n+1), n\}$. Fix $u \in I_\varepsilon$ and write $m^- = m_n^- = \lfloor \{h + \min(h, \varepsilon)\}(n+1) \rfloor - 1$; note $h + \min(h, \varepsilon) \leq \min(2h, h + \varepsilon)$. Four cases.

If neither endpoint is clipped, the interval has length $2h(n+1)$ and contains at least $\lfloor 2h(n+1) \rfloor - 1 \geq m^-$ integers.

If only the left endpoint is clipped, then $x(u) = 1$ and the count is $\lfloor y(u) \rfloor$. Either $y(u) = (u+h)(n+1) \geq (\varepsilon+h)(n+1)$, so the count is at least $\lfloor (h+\varepsilon)(n+1) \rfloor \geq m^- + 1$; or $y(u) = n$ and the count is $n \geq \lfloor (h+\varepsilon)(n+1) \rfloor - 1 \geq m^-$, using $h \leq 1 - \varepsilon$, i.e. $(h+\varepsilon)(n+1) \leq n+1$.

If only the right endpoint is clipped, then $y(u) = n$ and the count is $n - \lceil x(u) \rceil + 1 \geq n - x(u) \geq n - (1 - \varepsilon - h)(n+1) = (h+\varepsilon)(n+1) - 1$, hence at least $\lceil (h+\varepsilon)(n+1) \rceil - 1 \geq m^-$.

If both endpoints are clipped, the count is n , which is at least m^- as in the second case. In every case $M_j(u) \geq m^-$, and $k_j(u) = \lfloor \alpha M_j(u) \rfloor \geq \lfloor \alpha m_n^- \rfloor = K_n$. $\square$

Lemma A.9 (Finite-state structure). *For fixed j and fixed ranks, $u \mapsto \mathcal{W}_j(u)$ takes at most $S_n = 4n+1$ distinct values on $[0, 1]$.*

Proof. The membership indicator of observation i changes only at $\widehat{U}_{ij} \pm h$: at most $2n$ breakpoints, hence at most $2n+1$ constancy intervals plus the $2n$ breakpoint states. $\square$

Lemma A.10 (Local extreme count). *If $nh \geq 4$ and $\alpha m_n^- \geq 2$, then $m_n^- \geq nh/2$ and*

$$\frac{n\alpha h}{4} \leq K_n \leq 2n\alpha h + 1.$$

Proof. $m_n^- \geq \lfloor h(n+1) \rfloor - 1 \geq hn - 2 \geq nh/2$ for $nh \geq 4$, and $K_n = \lfloor \alpha m_n^- \rfloor \geq \alpha m_n^- / 2 \geq n\alpha h / 4$ when $\alpha m_n^- \geq 2$ (since $\lfloor x \rfloor \geq x/2$ for $x \geq 2$). For the upper bound, $m_n^- \leq 2h(n+1) - 1 \leq 2hn + 2h - 1 \leq 2hn + 1$, so $K_n \leq \alpha m_n^- \leq 2n\alpha h + \alpha \leq 2n\alpha h + 1$. $\square$

The rate conditions (7) force $nh \rightarrow \infty$ and $K_n \rightarrow \infty$, so Lemma A.10 applies for n large; it is the only conversion between K_n and $n\alpha h$ used in Appendix B.

A.4 Local Hill representation and uniform profile control

The scheme parallels the proof of Theorem 3 in Gardes and Podgorny [9] — oscillation control, exponential-spacings representation, slowly varying remainder — with the covering net over their compact matrix class replaced by exact union bounds over p coordinates and S_n window states.

Lemma A.11 (Conditional PIT and Rényi representation). *Fix j and condition on the column $\mathbf{U}_j = (U_{1j}, \dots, U_{nj})$. Under (E1):*

1. $V_{ij} = 1 - F_{Y|j}(Y_i | U_{ij})$, $i \leq n$, are i.i.d. standard uniform, and $Y_i = Q_j(1/V_{ij}, U_{ij})$ almost surely;
2. for any index set $\mathcal{W}$ of size m , deterministic given $\mathbf{U}_j$, with $V_{(1)} < \dots < V_{(m)}$ the ordered PIT values of its members and $k = \lfloor \alpha m \rfloor$,

$$L_k = \frac{1}{k} \sum_{r=1}^k \log \frac{V_{(k+1)}}{V_{(r)}}$$

is distributed as the mean of k i.i.d. standard exponentials, independent of $V_{(k+1)}$.

Proof. (i) Given $\mathbf{U}_j$, the pairs (U_{ij}, Y_i) are independent in i and Y_i has continuous conditional distribution function $F_{Y|j}(\cdot | U_{ij})$, so V_{ij} is standard uniform. For the representation, $Q_j(1/V_{ij}, U_{ij}) = F_{Y|j}^{-1}(F_{Y|j}(Y_i | U_{ij}) | U_{ij}) \leq Y_i$ always, with strict inequality only when Y_i lies in the interior of an interval on which $F_{Y|j}(\cdot | U_{ij})$ is constant; each such interval has conditional probability zero and there are countably many, so the representation holds almost surely. Strict monotonicity is not needed. Conditioning on the whole column, rather than on window counts as in Gardes and Podgorny [9, Appendix, Lemma 4], makes the window sets deterministic (Lemma A.8); only coordinate j is conditioned on, and no independence across coordinates is used.

(ii) With $W_i = -\log V_i$ standard exponentials and order statistics $W_{(1)} < \dots < W_{(m)}$, $W_{(m-r+1)} = -\log V_{(r)}$, Abel summation gives

$$kL_k = \sum_{r=1}^k \{W_{(m-r+1)} - W_{(m-k)}\} = \sum_{i=1}^k i \{W_{(m-i+1)} - W_{(m-i)}\},$$

and by the Rényi representation of exponential order statistics [16] the normalized top spacings $i\{W_{(m-i+1)} - W_{(m-i)}\}$, $i \leq k$, are i.i.d. standard exponential and independent of $W_{(m-k)}$, hence of $V_{(k+1)}$. $\square$

The next lemma is the reason no control of central quantiles is needed: only the $k+1$ largest responses enter a Hill statistic, and on the working event they are generated by PIT values below 3α .

Lemma A.12 (Tail-only freezing of the upper order statistics). *Fix $0 < \tau \leq 3\alpha < 1$ and a window state $\mathcal{W}$ of size m realized at u , with $k = \lfloor \alpha m \rfloor$. Write $Z_i = \log Q_j(V_{ij}^{-1}, U_{ij})$ and $\check{Z}_i = \log Q_j(V_{ij}^{-1}, u)$, and let $Z_{[1]} \geq \dots \geq Z_{[m]}$ and $\check{Z}_{[1]} \geq \dots \geq \check{Z}_{[m]}$ be the descending order statistics. On the event $\{\min_{i \in \mathcal{W}} V_{ij} \geq \tau\} \cap \{V_{(k+1)} \leq q := 3\alpha\}$, if every member satisfies $|U_{ij} - u| \leq r$, then*

$$\max_{1 \leq s \leq k+1} |Z_{[s]} - \check{Z}_{[s]}| \leq \omega := \omega_n^\uparrow(r, \tau, \alpha), \quad (\text{A.4})$$

and consequently the corresponding Hill statistics satisfy $|H - \check{H}| \leq 2\omega$.

Proof. Let $T = \{i \in \mathcal{W} : V_{ij} \leq q\}$. For $i \in T$ we have $V_{ij} \in [\tau, q]$, so (A.1) gives $|Z_i - \check{Z}_i| \leq \omega$. For $i \notin T$, monotonicity of $s \mapsto Q_j(s, \cdot)$ gives

$$Z_i \leq \log Q_j(q^{-1}, U_{ij}) \leq \log Q_j(q^{-1}, u) + \omega \leq \check{Z}_{[k+1]} + \omega, \quad (\text{A.5})$$

the last inequality because $V_{(k+1)} \leq q$ forces at least $k+1$ members to have $\check{Z}_i \geq \log Q_j(q^{-1}, u)$. Fix $s \leq k+1$. The indices carrying the s smallest PIT values lie in T , and their frozen values are at least $\check{Z}_{[s]}$ by monotonicity of $Q_j(\cdot, u)$; their actual values are therefore at least $\check{Z}_{[s]} - \omega$, whence $Z_{[s]} \geq \check{Z}_{[s]} - \omega$. Conversely at most $s-1$ indices of T have $\check{Z}_i > \check{Z}_{[s]}$; every remaining index of T has $Z_i \leq \check{Z}_{[s]} + \omega$, while (A.5) and $\check{Z}_{[k+1]} \leq \check{Z}_{[s]}$ give the same bound outside T . Hence $Z_{[s]} \leq \check{Z}_{[s]} + \omega$, proving (A.4). The Hill statistic is an average of differences of the first k order statistics and the $(k+1)$ st, so its perturbation is at most 2ω . $\square$

Lemma A.13 (Frozen-location Hill expansion). *Fix j , condition on $\mathbf{U}_j$, and fix a window state $\mathcal{W}$ of size m realized at $u \in I_\varepsilon$, with $k = \lfloor \alpha m \rfloor$ and Hill statistic H . On the event $D_n(\eta)$ intersected with $\{V_{ij} \geq \tau \ \forall i\}$ and $\{V_{(k+1)} \leq 3\alpha\}$,*

$$|H - \xi_j(u)| \leq 2\omega_n^\uparrow(h + \eta + \delta_n, \tau, \alpha) + \xi_j(u)|L_k - 1| + b_n(3\alpha)L_k. \quad (\text{A.6})$$

Proof. Spatial freezing. Each $i \in \mathcal{W}$ has $|\widehat{U}_{ij} - u| \leq h$, hence $|U_{ij} - u| \leq r_n := h + \eta + \delta_n$ on $D_n(\eta)$. Lemma A.12, applied with $r = r_n$, gives $|H - \check{H}| \leq 2\omega_n^\uparrow(r_n, \tau, \alpha)$, where $\check{H}$ is the Hill statistic of the frozen responses $\check{Y}_i = Q_j(1/V_{ij}, u)$.

Exact decomposition. Since $s \mapsto Q_j(s, u)$ is nondecreasing, $\check{Y}_{(m-r+1)} = Q_j(1/V_{(r)}, u)$, and $\log Q_j(s, u) = \xi_j(u) \log s + \log \ell_j(s, u)$ yields

$$\check{H} = \xi_j(u)L_k + R, \quad R = \frac{1}{k} \sum_{r=1}^k \log \frac{\ell_j(1/V_{(r)}, u)}{\ell_j(1/V_{(k+1)}, u)}.$$

Remainder. On $\{V_{(k+1)} \leq 3\alpha\}$, each ratio is $\ell_j(ts, u)/\ell_j(s, u)$ with $s = 1/V_{(k+1)} \geq (3\alpha)^{-1}$ and $t = V_{(k+1)}/V_{(r)} > 1$, so $|\log(\cdot)| \leq b_n(3\alpha) \log(V_{(k+1)}/V_{(r)})$; averaging gives $|R| \leq b_n(3\alpha)L_k$. $\square$

The decomposition behind (A.6) is the shape of the whole argument:

$$\widehat{\xi}_j(u) - \xi_j(u) = \text{rank/spatial error} + \text{exponential fluctuation} + \text{tail remainder}.$$

Proposition A.14 (Finite-sample uniform profile bound). *Assume (C1)–(C2), (S), the structural part of (E1) (continuity of every $F_{Y|j}(\cdot | u)$), hence the quantile representation), that $b_n(3\alpha)$ and $\omega_n^\uparrow(\cdot, \cdot, \cdot)$ are finite, $0 < h \leq 1 - \varepsilon$, $0 < \alpha \leq 1/3$, and $K_n \geq 2$. For every $\eta > 0$, $0 < t \leq 1$, $0 < \tau \leq 3\alpha$,*

$$\begin{aligned} \mathbb{P} \left\{ \max_{j \leq p} \sup_{u \in I_\varepsilon} |\widehat{\xi}_j(u) - \xi_j(u)| > 2\omega_n^\uparrow(h + \eta + \delta_n, \tau, \alpha) + \Gamma t + (1+t)b_n(3\alpha) \right\} \\ \leq \underbrace{2pe^{-2n\eta^2}}_{\text{ranks}} + \underbrace{pn\tau}_{\text{PIT truncation}} + \underbrace{2pS_n e^{-K_n t^2/4}}_{\text{spacings}} + \underbrace{pS_n e^{-(K_n+1)/4}}_{\text{threshold}}. \end{aligned} \quad (\text{A.7})$$

Proof. Call a window state *admissible* for coordinate j if it equals $\mathcal{W}_j(u)$ for some $u \in I_\varepsilon$; by Lemmas A.8 and A.9 there are at most S_n admissible states per coordinate and each has $k \geq K_n \geq 2$, so L_k and $V_{(k+1)}$ are well defined on every one of them. (States realized only outside I_ε may have smaller windows; they are not needed.) Define the exceptional events $T_1 = D_n(\eta)^c$; $T_2 = \{\exists i, j :$

$V_{ij} < \tau\}$; $T_3 = \{\exists j, \text{ admissible state } \mathcal{W} : |L_k - 1| > t\}$; $T_4 = \{\exists j, \text{ admissible state } \mathcal{W} : V_{(k+1)} > 3\alpha\}$.

Outside $T_1 \cup T_2 \cup T_3 \cup T_4$, fix j and $u \in I_\varepsilon$: the realized window is an admissible state, and Lemma A.13 with $\xi_j \leq \Gamma$, $|L_k - 1| \leq t$, $L_k \leq 1 + t$ gives the displayed deviation bound; the modulus $\omega_n^\uparrow(h + \eta + \delta_n, \tau, \alpha)$ covers every selected observation because its definition allows the displaced point $v = U_{ij}$ to range over all of $[0, 1]$, and, by Lemma A.12, only PIT values in $[\tau, 3\alpha]$ are involved. The upper truncation $V_{ij} \leq 1 - \tau$ is never used.

Probabilities: $\mathbb{P}(T_1) \leq 2pe^{-2n\eta^2}$ is Lemma A.7; $\mathbb{P}(T_2) \leq pn\tau$ since each V_{ij} is uniform. For T_3, T_4 , condition on $\mathbf{U}_j$; the admissible states are then deterministic with $k \geq K_n$, Lemma A.11 applies to each, and Lemmas A.2 and A.3 give conditional bounds $2e^{-kt^2/4} \leq 2e^{-K_nt^2/4}$ and $e^{-(k+1)/4} \leq e^{-(K_n+1)/4}$, both decreasing in k . A union bound over the at most pS_n admissible states and integration over $\mathbf{U}_j$ complete the proof; overlapping states and dependent coordinates cost nothing beyond the union bound. $\square$

A.5 From profiles to scores

Recall the score grid $G = \{r\delta_n : r = 1, \dots, n\} \cap I_\varepsilon$, of size $L_n = |G|$, and that $K_n \geq 2$ makes every grid point usable, so $\widehat{\Psi}_j = L_n^{-1} \sum_{u \in G} \widehat{\xi}_j(u)$.

Lemma A.15 (Quadrature). *If $L_n \geq 1$, then for every j ,*

$$\left| L_n^{-1} \sum_{u \in G} \widehat{\xi}_j(u) - \Psi_j \right| \leq q_n := \omega_\xi(\delta_n) + \frac{4\Gamma\delta_n}{1 - 2\varepsilon}. \quad (\text{A.8})$$

Proof. Write the grid points $u_1 < \dots < u_{L_n}$, consecutive multiples of δ_n . If $L_n = 1$, then I_ε contains a single multiple of δ_n , which forces $|I_\varepsilon| < 2\delta_n$; the left side is at most Γ , while $4\Gamma\delta_n/|I_\varepsilon| > 2\Gamma$, so (A.8) holds trivially. Assume $L_n \geq 2$ and partition I_ε exactly: $C_1 = [\varepsilon, u_1 + \delta_n/2]$, $C_{L_n} = [u_{L_n} - \delta_n/2, 1 - \varepsilon]$, and $C_r = [u_r - \delta_n/2, u_r + \delta_n/2]$ otherwise. Interior cells have length δ_n , boundary cells length in $[\delta_n/2, 3\delta_n/2]$, and every point of C_r is within δ_n of u_r . Then

$$\left| \frac{1}{L_n} \sum_r \widehat{\xi}_j(u_r) - \frac{1}{|I_\varepsilon|} \int_{I_\varepsilon} \widehat{\xi}_j \right| \leq \frac{1}{|I_\varepsilon|} \sum_r \int_{C_r} |\widehat{\xi}_j - \widehat{\xi}_j(u_r)| + \Gamma \sum_r \left| \frac{\lambda(C_r)}{|I_\varepsilon|} - \frac{1}{L_n} \right|.$$

The first term is at most $\omega_\xi(\delta_n)$. For the second, both weight vectors sum to one; $|I_\varepsilon| = \sum_r \lambda(C_r) \in [(L_n - 1)\delta_n, (L_n + 1)\delta_n]$ gives $||I_\varepsilon|/L_n - \delta_n| \leq \delta_n/L_n$, so interior cells contribute at most δ_n in total and the two boundary cells at most $\delta_n/2 + \delta_n/L_n \leq \delta_n$ each; the sum is at most $3\delta_n$, and the second term is at most $3\Gamma\delta_n/|I_\varepsilon| \leq 4\Gamma\delta_n/|I_\varepsilon|$. $\square$

Proposition A.16 (Finite-sample score bound). *Under the conditions of Proposition A.14 (with in addition $\omega_\xi(\delta_n) < \infty$) and $L_n \geq 1$, for every $\eta > 0$, $0 < t \leq 1$, $0 < \tau < 1/2$,*

$$\mathbb{P}\left\{ \max_{j \leq p} |\widehat{\Psi}_j - \Psi_j| > B_n(\eta, t, \tau) + q_n \right\} \leq 2pe^{-2n\eta^2} + pn\tau + 2pS_ne^{-K_nt^2/4} + pS_ne^{-(K_n+1)/4}, \quad (\text{A.9})$$

where $B_n(\eta, t, \tau) = 2\omega_n^\uparrow(h + \eta + \delta_n, \tau, \alpha) + \Gamma t + (1 + t)b_n(3\alpha)$.

Proof. Pathwise,

$$|\widehat{\Psi}_j - \Psi_j| \leq \sup_{u \in I_\varepsilon} |\widehat{\xi}_j(u) - \xi_j(u)| + \left| L_n^{-1} \sum_{u \in G} \widehat{\xi}_j(u) - \Psi_j \right|,$$

and Lemma A.15 bounds the second term by q_n on every sample path; reusing the same observations for estimation and averaging therefore costs no extra probability. Proposition A.14 gives (A.9). $\square$

Lemma A.17 (Deterministic score separation). *Suppose $\Delta_{\min} > 0$ and, on some event,*

$$\max_{j \leq p} |\hat{\Psi}_j - \Psi_j| < \frac{\Delta_{\min}}{2}.$$

Then, on that event, every active coordinate has a strictly smaller score than every inactive coordinate; consequently $A \subseteq \hat{A}_d$ for every $d \geq s$, and $\hat{A}_s = A$.

Proof. Inactive coordinates have $\Psi_j = \gamma^*$ (Proposition 2.2, proved independently of this lemma in Appendix B), hence $\hat{\Psi}_j > \gamma^* - \Delta_{\min}/2$; active coordinates have $\Psi_j \leq \gamma^* - \Delta_{\min}$, hence $\hat{\Psi}_j < \gamma^* - \Delta_{\min}/2$. Every active score therefore precedes every inactive score in the increasing order, and the screening and recovery statements follow from the definition of $\hat{A}_d$. $\square$

B Proofs of the main results

This appendix proves the propositions of Section 2 and the theorems of Section 3 from the auxiliary results of Appendix A.

B.1 Proof of Proposition 2.1

For fixed $t > 0$, with N and K_y^* as in Lemmas A.4 and A.5 applied to $\mathcal{G}$,

$$\frac{N(ty)}{N(y)} = \int_{[\underline{\gamma}, \Gamma]} t^{-1/g} K_y^*(dg) \longrightarrow t^{-1/\xi}$$

almost surely, by Lemma A.5 and boundedness and continuity of $g \mapsto t^{-1/g}$ on $[\underline{\gamma}, \Gamma]$. Lemma A.4 transfers the limit to $\mathbb{P}(Y > ty \mid \mathcal{G})/\mathbb{P}(Y > y \mid \mathcal{G})$. Convergence simultaneously for all $t > 0$ follows from rational t , monotonicity of the ratio in t , and continuity of $t \mapsto t^{-1/\xi}$.

For the fibre identity (3), fix $u \in I_\varepsilon$ and apply the above on the probability space $([0, 1]^{p-1}, K_j(u, \cdot))$ with the trivial σ -field; $\xi_j(u)$ is by definition the essential supremum of $v \mapsto \gamma(u, v)$ under $K_j(u, \cdot)$. Let $m = \max_v \gamma(u, v)$, attained by compactness and continuity. The essential supremum is at most m . Conversely, for $\eta > 0$ the set $\{v : \gamma(u, v) > m - \eta\}$ is open and nonempty, so has positive $K_j(u, \cdot)$ -probability under (S); the essential supremum is at least $m - \eta$ for every $\eta > 0$. $\square$

B.2 Proof of Proposition 2.2

Claim 1. For $j \notin A$, $\gamma(u, v) = g\{(u, v)_A\}$ does not involve u , and as v ranges over the fibre the active subvector ranges over all of $[0, 1]^s$; by (3), $\xi_j(u) = \max_w g(w) = \gamma^*$ for every u .

Claim 2. By (3), $\xi_j(u) = \gamma^*$ exactly when the fibre through u intersects $\mathcal{M}$, i.e. when $u \in \pi_j(\mathcal{M})$.

Claim 3. $u \mapsto \gamma^* - \xi_j(u)$ is continuous (Lemma A.6), nonnegative, and by Claim 2 vanishes on I_ε exactly on $I_\varepsilon \cap \pi_j(\mathcal{M})$. A continuous nonnegative function has positive integral over I_ε if and only if its positivity set has positive Lebesgue measure.

Final statement. Write $\mathcal{M}^A$ for the maximizer set of g in $[0, 1]^s$; then $\mathcal{M} = \{u : u_A \in \mathcal{M}^A\}$ and, for $j \in A$, $\pi_j(\mathcal{M})$ is the projection of $\mathcal{M}^A$, a finite set when $\mathcal{M}^A$ is finite; its intersection with I_ε is Lebesgue-null and Claim 3 applies. $\square$

B.3 Proof of Proposition 2.3

Every x on the fibre $x_j = u$ has $\|x_A - u_A^*\| \geq |u - u_j^*|$, so every point of the fibre satisfies $\gamma(x) \leq \gamma^* - c_1 |u - u_j^*|^\nu$; taking the fibre maximum gives $\xi_j(u) \leq \gamma^* - c_1 |u - u_j^*|^\nu$. Averaging over $u \in I_\varepsilon$ gives the score bound. For $\nu = 2$, $\mathbb{E}(W - u_j^*)^2 = \text{Var}(W) + (1/2 - u_j^*)^2 = (1 - 2\varepsilon)^2/12 + (1/2 - u_j^*)^2$. $\square$

B.4 Proof of Theorem 3.1

Choose $\eta_n = \sqrt{\log(pn)/n}$, $\tau = \tau_n = (pn)^{-2}$, and $t_n = \sqrt{(4/K_n) \log\{pS_n(pn)\}}$. The four probability terms in (A.9) are then each $o(1)$: the first is $2pe^{-2\log(pn)} = 2p(pn)^{-2} \rightarrow 0$; the second is $2pn(pn)^{-2} \rightarrow 0$; the third is $2pS_n\{pS_n(pn)\}^{-1} = 2(pn)^{-1} \rightarrow 0$; the fourth is $pS_n e^{-(K_n+1)/4} \rightarrow 0$ because $\log(pS_n) \leq C \log(pn) = o(K_n)$ by the second condition in (7) and Lemma A.10. The same condition makes $t_n \leq 1$ for large n , as required by Lemma A.2, and $\Gamma t_n = O(\sqrt{\log(pn)/(n\alpha h)})$, again by Lemma A.10.

It remains to check that the bias part of $B_n(\eta_n, t_n, \tau_n) + q_n$ in Proposition A.16 is $o(\Delta_{\min})$. First, $\log(pn) = o(nh^2)$ gives $\eta_n + \delta_n = o(h)$, so $h + \eta_n + \delta_n \leq C_0 h$ eventually and $2\omega_n^\uparrow(h + \eta_n + \delta_n, \tau_n, \alpha) \leq 2\omega_n^\uparrow(C_0 h, \tau_n, \alpha) = o(\Delta_{\min})$ by (E3) and monotonicity of $\omega_n^\uparrow(\cdot, \tau, \alpha)$. Second, $(1 + t_n)b_n(3\alpha) \leq 2b_n(3\alpha) = o(\Delta_{\min})$ by (E2). Third, $q_n = \omega_\xi(\delta_n) + O(1/n)$: the first part is $o(\Delta_{\min})$ by (E1), and $1/n \leq \sqrt{\log(pn)/(n\alpha h)}$ since $\alpha h \leq 1 \leq n \log(pn)$, so the second part is absorbed into the stochastic rate. Collecting terms, outside a vanishing probability,

$$\max_{j \leq p} |\hat{\Psi}_j - \Psi_j| \leq a_n + \beta_n, \quad a_n = O\left(\sqrt{\log(pn)/(n\alpha h)}\right), \quad \beta_n = o(\Delta_{\min}),$$

which is the assertion. $\square$

B.5 Proof of Theorem 3.2 and Corollary 3.3

Under (8), $\sqrt{\log(pn)/(n\alpha h)} = o(\Delta_{\min})$, so the bound of Theorem 3.1 is below $\Delta_{\min}/2$ for n large, outside a vanishing probability. On the complementary event, Lemma A.17 places every active score strictly before every inactive score, so $A \subseteq \hat{A}_d$ for every $d \geq s$, and $\hat{A}_s = A$ when $d = s$. $\square$

C Primitive sufficient conditions

Conditions (E1)–(E3) are stated on the projected quantiles. This appendix gives primitive sufficient conditions — conditions on γ , on the full conditional survival function, and on the fibre kernels $K_j(u, \cdot)$ — that imply (E1), (E2) and the tail-only condition (E3 $^\uparrow$) of Appendix A, with explicit rates, and verifies them for the simulation models of Section 4.

All constants and envelopes below are uniform in the triangular array, in $j \leq p_n$, and on the enlarged interval $I_{\varepsilon/2} = [\varepsilon/2, 1 - \varepsilon/2]$; scores continue to be integrated over I_ε . For $x = \iota_j(u, v)$ put

$$\xi_j(u) = \text{ess sup}_{v \sim K_j(u, \cdot)} \gamma\{\iota_j(u, v)\}, \quad D_{j,u}(v) = \frac{1}{\gamma\{\iota_j(u, v)\}} - \frac{1}{\xi_j(u)} \geq 0. \quad (\text{C.1})$$

The reciprocal-index deficit $D_{j,u}$, rather than $\xi_j(u) - \gamma$, is what makes the Laplace representation below exact.

C.1 Why (C1)–(C2) do not suffice

Proposition C.1 (Necessary repairs). *The assumptions (C1)–(C2), even together with (S), imply none of (E1), (E2), (E3).*

Proof. (C1) constrains only the upper tail and does not exclude atoms, so it does not imply the continuity of the projected conditional distribution required by (E1).

For (E2), a level bound on a Hall remainder is insufficient. Let $\bar{F}(e^t) = \exp\{-t + r(t)\}$, where r is a sum of disjoint smooth negative bumps centred at $t_m \rightarrow \infty$, the m th of height $a_m = e^{-\beta t_m}$, with

a decreasing side of width $o(a_m)$ and a recovery side of slope at most $1/2$. Then $-1 + r'(t) < 0$, so $\bar{F}$ is a valid survival function and $|r(t)| \leq e^{-\beta t}$, yet $r'(t_m) \rightarrow -\infty$. With $T(x) = \log Q(e^x)$, quantile inversion gives

$$\frac{d}{dx} \log\{Q(e^x)e^{-x}\} = \frac{r'\{T(x)\}}{1 - r'\{T(x)\}},$$

whose absolute value stays away from zero along a subsequence, so the normalized increment in (E2) need not vanish.

Finally, make all conditional survival functions identical above a fixed level but let the conditional median below it vary discontinuously with the conditioning coordinate. This changes none of (C1), (C2), (S), and none of the tail quantities, yet the printed (E3) fails because its range $w \in [\tau_n, 1 - \tau_n]$ contains $w = 1/2$. This last obstruction is an artefact of the statement, not of the method: by Lemma A.12 the proofs use only $w \leq 3\alpha_n$, so the correct repair is to weaken (E3) to (E3[†]) rather than to impose regularity on the body of the distribution. $\square$

C.2 The primitive conditions

(P1) *Geometry and atomlessness.* For constants $0 < \underline{\gamma} \leq \Gamma < \infty$ and $L_\gamma < \infty$,

$$\underline{\gamma} \leq \gamma(x) \leq \Gamma, \quad |\gamma(x) - \gamma(x')| \leq L_\gamma \|x - x'\|_\infty.$$

Full fibre support (S) holds, and every conditional law $Y \mid U = x$ is atomless.

(P2) *Differentiable uniform tail approximation.* For $t \geq t_0$,

$$\mathbb{P}(Y > e^t \mid U = x) = c(x)e^{-t/\gamma(x)}\{1 + r(t, x)\}, \quad 0 < c_- \leq c(x) \leq c_+ < \infty,$$

with $\sup_x |r(t, x)| \leq R_0(t)$ and $\sup_x |\partial_t r(t, x)| \leq R_1(t)$, where $R_0(t) \rightarrow 0$ and $\bar{R}_1(T) = \sup_{t \geq T} R_1(t) \rightarrow 0$; eventually $R_0(t) \leq 1/2$. The maps $r(\cdot, x)$ are locally absolutely continuous.

(P3) *Uniform doubling of the fibre-deficit mass.* With

$$H_{j,u}(z) = \int \mathbf{1}_{\{D_{j,u}(v) \leq z\}} c\{\iota_j(u, v)\} K_j(u, dv),$$

there are $z_0, h_0 > 0$ and $C_D < \infty$ with

$$H_{j,u}(z_0) \geq h_0, \quad H_{j,u}(2z) \leq C_D H_{j,u}(z) \quad (0 < z \leq z_0/2)$$

uniformly. This allows an atom at $D = 0$, includes the degenerate case $D \equiv 0$ of an inactive coordinate, and rules out only fibres on which the probability of coming near the maximal index becomes arbitrarily thinner under a fixed rescaling.

(P4) *Horizontal coupling of the primitive tail.* Whenever $u \in I_\varepsilon$, $v \in I_{\varepsilon/2}$ and $|u - v| \leq \varepsilon/2$, the kernels $K_j(u, \cdot)$ and $K_j(v, \cdot)$ admit a coupling $X_u = \iota_j(u, V_u)$, $X_v = \iota_j(v, V_v)$ such that, almost surely and for every $t \geq t_0$,

$$\left| \frac{1}{\gamma(X_v)} - \frac{1}{\gamma(X_u)} \right| \leq L_D |v - u|,$$

$$|\log c(X_v) - \log c(X_u)| + |\log\{1 + r(t, X_v)\} - \log\{1 + r(t, X_u)\}| \leq L_H |v - u| + \zeta_n,$$

with $\zeta_n \geq 0$ and $\zeta_n = o(\Delta_{\min})$. This concerns only the primitive conditional tail of (P2), never the body.

A sufficient construction for (P4) is a coupling with $\|X_v - X_u\|_\infty \leq C_K |v - u|$ together with Lipschitz $1/\gamma$, $\log c$ and $\log(1 + r)$ on the coupled states; (P4) is weaker because it records only the three quantities actually used.

C.3 The transfer theorem

Lemma C.2 (Tilted deficit bound). *Under (P3), for all sufficiently large t the quantity*

$$B_{j,u}(t) = \int e^{-tD_{j,u}(v)} c\{\iota_j(u, v)\} K_j(u, dv)$$

satisfies $B_{j,u}(t/2) \leq C_B B_{j,u}(t)$ for a uniform constant C_B , and consequently

$$\frac{\int D_{j,u}(v) e^{-tD_{j,u}(v)} c\{\iota_j(u, v)\} K_j(u, dv)}{B_{j,u}(t)} \leq \frac{2C_B}{et} \quad (\text{C.2})$$

uniformly in n, j, u .

Proof. Suppress (j, u) . For $t \geq z_0^{-1}$, restricting the integral to $\{D \leq 1/t\}$ gives $B(t) \geq e^{-1} H(1/t)$. Split $B(t/2)$ over $[0, 1/t]$ and the dyadic bands $(2^m/t, 2^{m+1}/t]$. Before the bands reach z_0 their total contribution is at most

$$H(1/t) \left[1 + \sum_{m \geq 0} e^{-2^{m-1}} C_D^{m+1} \right],$$

a convergent series. Iterating the doubling bound backwards from z_0 gives $H(1/t) \geq ct^{-\log_2 C_D}$ for a uniform $c > 0$, so the remaining part, which starts no lower than $z_0/2$ and is at most $c_+ e^{-tz_0/4}$, is also bounded by a constant times $H(1/t)$. Combining the two displays proves the doubling property of B . For $z \geq 0$, $ze^{-tz} \leq (2/e)t^{-1}e^{-tz/2}$; integrating and using the doubling property gives (C.2). $\square$

Theorem C.3 (Primitive sufficient conditions). *Assume (P1)–(P4), let $\alpha_n \rightarrow 0$, and suppose eventually $C_0 h_n \leq \varepsilon/2$ and $\tau_n = (p_n n)^{-2} \leq 3\alpha_n < 1$. Put $x_\alpha = \log\{(3\alpha_n)^{-1}\}$ and $c_\gamma = \underline{\gamma}/2$. Then, after a common deterministic onset,*

$$\omega_\xi(r) \leq L_\gamma r, \quad (\text{C.3})$$

$$b_n(3\alpha_n) \leq \frac{C_E}{x_\alpha} + 4\Gamma^2 \bar{R}_1(c_\gamma x_\alpha), \quad (\text{C.4})$$

$$\omega_n^\uparrow(C_0 h_n, \tau_n, \alpha_n) \leq 2\Gamma C_0 h_n [L_H + L_D \Gamma \{2 \log(p_n n) + C_Q\}] + 2\Gamma \zeta_n, \quad (\text{C.5})$$

where $C_E < \infty$ depends only on $\underline{\gamma}$, Γ , C_B and the onset constants, and $C_Q = \max\{0, \log(2c_+)\}$. Consequently (E1), (E2) and (E3[†]) hold as soon as

$$n^{-1} = o(\Delta_{\min}), \quad x_\alpha^{-1} + \bar{R}_1(c_\gamma x_\alpha) = o(\Delta_{\min}), \quad h_n \log(p_n n) + \zeta_n = o(\Delta_{\min}). \quad (\text{C.6})$$

For a fixed gap and a Hall derivative envelope of order $O(t^{-1})$ or smaller, the substantive requirement from (E2) is simply $\log(1/\alpha_n) \rightarrow \infty$.

Proof. Step 1: (E1). Under (S), $\xi_j(u) = \max_v \gamma\{\iota_j(u, v)\}$; comparing the two maxima at the same v in both directions gives (C.3), hence $\omega_\xi(1/n) \leq L_\gamma/n$. A mixture of atomless laws is atomless, because $\mathbb{P}(Y = y \mid U_j = u) = \int \mathbb{P}\{Y = y \mid U = \iota_j(u, v)\} K_j(u, dv) = 0$; this gives the continuity and the conditional quantile representation required by (E1).

Step 2: exact projected factorization. All derivative identities below hold almost everywhere, and local absolute continuity justifies their integrated forms. Let

$$\mathcal{A}_{j,u}(t) = \int e^{-tD_{j,u}(v)} c\{\iota_j(u, v)\} \{1 + r(t, \iota_j(u, v))\} K_j(u, dv),$$

so that $\bar{F}_j(e^t \mid u) = e^{-t/\xi_j(u)} \mathcal{A}_{j,u}(t)$ exactly. With $M_{j,u}(t) = \int D_{j,u} e^{-tD_{j,u}c} K_j(u, dv)$, (P2), Lemma C.2 and $R_0 \leq 1/2$ give

$$\left| \frac{\mathcal{A}'_{j,u}(t)}{\mathcal{A}_{j,u}(t)} \right| \leq \frac{\{1 + R_0(t)\}M_{j,u}(t) + R_1(t)B_{j,u}(t)}{\{1 - R_0(t)\}B_{j,u}(t)} \leq \frac{C_A}{t} + 2R_1(t) =: g(t), \quad (\text{C.7})$$

and $g(t) \rightarrow 0$. Moreover $(1 - R_0)B_{j,u} \leq \mathcal{A}_{j,u} \leq (1 + R_0)B_{j,u}$, while the same tilted bound gives $-(\log B_{j,u})'(t) \leq C/t$. Integrating from a fixed t_1 at which $R_0(t_1) \leq 1/2$, and using $D \leq \underline{\gamma}^{-1}$ and $c \geq c_-$ to bound $B_{j,u}(t_1)$ away from zero, yields uniform positive constants with $ct^{-C} \leq \mathcal{A}_{j,u}(t) \leq 2c_+$, hence

$$|\log \mathcal{A}_{j,u}(t)| \leq C \log t \quad (\text{C.8})$$

uniformly for large t .

Step 3: inversion and (E2). Let $T_{j,u}(x) = \log Q_j(e^x, u)$. Once $g(t) < 1/(2\Gamma)$ the map $t \mapsto \bar{F}_j(e^t \mid u)$ is strictly decreasing and invertible, and

$$x = \frac{T_{j,u}(x)}{\xi_j(u)} - \log \mathcal{A}_{j,u}\{T_{j,u}(x)\}. \quad (\text{C.9})$$

By (C.8), $T_{j,u}(x) \geq c_\gamma x$ for all sufficiently large x . Writing $Q_j(e^x, u) = e^{\xi_j(u)x} \ell_j(e^x, u)$ and differentiating (C.9),

$$\frac{d}{dx} \log \ell_j(e^x, u) = \frac{\xi_j(u)^2 a_{j,u}\{T_{j,u}(x)\}}{1 - \xi_j(u) a_{j,u}\{T_{j,u}(x)\}}, \quad a_{j,u} = \frac{\mathcal{A}'_{j,u}}{\mathcal{A}_{j,u}},$$

so (C.7) and $T_{j,u}(x) \geq c_\gamma x$ give

$$\sup_{j,u} \left| \frac{d}{dx} \log \ell_j(e^x, u) \right| \leq \frac{C_E}{x} + 4\Gamma^2 \bar{R}_1(c_\gamma x).$$

For $q > 1$ the ratio appearing in (E2) is an average of this derivative over $[\log s, \log(qs)]$; taking $s \geq (3\alpha_n)^{-1}$ and then all suprema proves (C.4).

Step 4: tail coupling and (E3[†]). For $t \geq t_0$ set $f_t(x) = c(x)e^{-t/\gamma(x)}\{1 + r(t, x)\}$. Under the coupling of (P4), with $d = |u - v|$,

$$e^{-\delta_t} f_t(X_u) \leq f_t(X_v) \leq e^{\delta_t} f_t(X_u), \quad \delta_t = (L_H + tL_D)d + \zeta_n,$$

so taking expectations gives $|\log \bar{F}_j(e^t \mid v) - \log \bar{F}_j(e^t \mid u)| \leq \delta_t$. By (C.7), after increasing the onset if necessary, $-\frac{d}{dt} \log \bar{F}_j(e^t \mid u) \geq 1/(2\Gamma)$, while $\mathcal{A}_{j,u} \leq 2c_+$ and (C.9) give $T_{j,u}(x) \leq \Gamma(x + C_Q)$. Evaluating the survival comparison at $t = T_{j,u}(x)$ and using the inverse-slope bound for the v -law,

$$|T_{j,v}(x) - T_{j,u}(x)| \leq 2\Gamma[L_H + L_D\Gamma(x + C_Q)]d + \zeta_n.$$

Taking $x = \log(1/w)$ with $w \in [\tau_n, 3\alpha_n]$: the lower endpoint $x \geq x_\alpha \rightarrow \infty$ places both quantiles beyond the common onset, while $x \leq 2\log(p_n n)$; for $d \leq C_0 h_n \leq \varepsilon/2$ the displaced point lies in $I_{\varepsilon/2}$. Taking suprema proves (C.5). $\square$

Remark C.4 (The body is never used). Step 4 controls only $w \leq 3\alpha_n$. This is exactly the range that Lemma A.12 shows to be relevant, and it is why (P4) constrains the primitive tail alone. No coupling of central quantiles is required anywhere.

C.4 Verifying deficit-mass doubling

Proposition C.5 (A one-step ratio bound implies (P3)). *Suppose that, uniformly in n, j, u , either $H_{j,u}(0) \geq h_0 > 0$, or*

$$H_{j,u}(z) = a_{j,u} z^{\kappa_{j,n}} L_{j,u}(1/z) \{1 + o(1)\}, \quad z \downarrow 0, \quad (\text{C.10})$$

with $a_{j,u}$ bounded above and away from zero, $0 \leq \kappa_{j,n} \leq \kappa_+ < \infty$, the $o(1)$ uniform, $L_{j,u}$ bounded above and away from zero at one common point beyond the uniform onset, and

$$\sup_{n,j,u} \frac{L_{j,u}(t/2)}{L_{j,u}(t)} \leq C_L \quad (\text{C.11})$$

for all sufficiently large t . Then (P3) holds.

Proof. In the atomic case $H(z) \geq h_0$ for every $z > 0$, so $H(2z) \leq c_+ \leq (c_+/h_0)H(z)$. Otherwise

$$\frac{H_{j,u}(2z)}{H_{j,u}(z)} = 2^{\kappa_{j,n}} \frac{L_{j,u}\{1/(2z)\}}{L_{j,u}(1/z)} \{1 + o(1)\},$$

uniformly bounded by (C.11) and $\kappa_{j,n} \leq \kappa_+$. Taking a fixed small z_0 , the expansion and the uniform lower bound on $a_{j,u} L_{j,u}(1/z_0)$ give $H_{j,u}(z_0) \geq h'_0 > 0$. $\square$

A pure two-sided power margin is not necessary, and is in fact false for the models below: their corner margins carry powers of $\log(1/z)$ and factors $\exp\{b\sqrt{\log(1/z)}\}$, which are slowly varying and therefore satisfy (C.11).

C.5 Verification for A1–A3 and B1

The active set in all four simulation models is $\{1, 2, 3, 4\}$ and the fibre maximum is attained when every free active PIT coordinate equals zero. The following boundary calculation is the only model-specific ingredient; its dimension is at most four, independently of p_n .

Lemma C.6 (Gaussian corner). *Let $X \sim N_d(\mu, \Sigma)$ with $1 \leq d \leq d_0$, let $\Omega = \Sigma^{-1}$, and let the parameters range over a compact family of positive definite covariance matrices and bounded means with $\min_i (\Omega \mathbf{1})_i \geq \nu_0 > 0$. Suppose that on a fixed neighbourhood N_0 of the lower corner D is uniformly C^2 with*

$$D\{\Phi(X)\} = \sum_{i=1}^d a_i \Phi(X_i) + O(\|\Phi(X)\|^2), \quad 0 < a_- \leq a_i \leq a_+ < \infty,$$

and that $c_ \sum_i u_i \leq D(u) \leq C_* \sum_i u_i$ on N_0 with $\inf_{u \notin N_0} D(u) > 0$. Write $\kappa = \mathbf{1}^\top \Omega \mathbf{1}$, $\nu = \Omega \mathbf{1}$, $b = \mathbf{1}^\top \Omega \mu$ and $r_z = -\Phi^{-1}(z)$. Then, uniformly over the family,*

$$\mathbb{P}[D\{\Phi(X)\} \leq z] = C(\mu, \Sigma, a) z^\kappa \{\sqrt{2\pi} r_z\}^{\kappa-d} e^{-br_z} \{1 + O(r_z^{-1})\}, \quad (\text{C.12})$$

with $C(\mu, \Sigma, a) = |\Sigma|^{-1/2} e^{-\mu^\top \Omega \mu / 2} \prod_i \Gamma(\nu_i) a_i^{-\nu_i} / \Gamma(\kappa + 1)$. The conclusion persists after multiplication by a positive weight that is uniformly Lipschitz at the corner and bounded there above and away from zero. In particular (P3) holds.

Proof. The density of $U = \Phi(X)$ is

$$f_U(u) = |\Sigma|^{-1/2} \exp\left\{-\frac{1}{2}z(u)^\top (\Omega - I)z(u) + z(u)^\top \Omega \mu - \frac{1}{2}\mu^\top \Omega \mu\right\}, \quad z_i(u) = \Phi^{-1}(u_i).$$

Set $u_i = \delta x_i$ and $r = r_\delta$. Uniformly on compact subsets of $(0, \infty)^d$, $\Phi^{-1}(\delta x_i) = -r + r^{-1} \log x_i + o(r^{-1})$ and $e^{-r^2/2} = \sqrt{2\pi r} \delta \{1 + O(r^{-2})\}$, so with $\nu = \Omega \mathbf{1}$

$$f_U(\delta x) \delta^d = |\Sigma|^{-1/2} e^{-\mu^\top \Omega \mu / 2} \delta^\kappa \{\sqrt{2\pi r}\}^{\kappa-d} e^{-br} \prod_i x_i^{\nu_i-1} \{1 + o(1)\}.$$

The two-sided bound on D confines $\{D(\delta x) \leq \delta\}$ to a fixed rescaled simplex and excludes points outside N_0 . Choose $0 < \theta < \frac{1}{2} \min\{1, \nu_0, \underline{\lambda}\}$ with $\underline{\lambda} = \inf \lambda_{\min}(\Omega)$, and split the simplex into $\mathcal{B}_r = \{x : \min_i x_i < e^{-\sqrt{r}}\}$ and its complement. On $\mathcal{B}_r$, Mills bounds and convexity of the Gaussian quadratic form give the envelope $C \prod_i x_i^{\theta-1}$, and a one-dimensional beta integral bounds the normalized mass of $\mathcal{B}_r$ by $C e^{-\theta\sqrt{r}} = o(r^{-1})$. On $\mathcal{B}_r^c$ we have $|\log x_i| \leq \sqrt{r}$ from below and a uniform bound from above, and the uniform normal-quantile expansion

$$\Phi^{-1}(\delta x_i) = -r + \frac{\log x_i}{r} + O\left(\frac{1 + |\log x_i| + (\log x_i)^2 + \log r}{r^3}\right)$$

shows that the quadratic Gaussian term contributes a relative error r^{-1} times a fixed polynomial in $1 + \sum_i |\log x_i|$, whose Dirichlet log moments are finite uniformly because $\nu_i \geq \nu_0$. The quadratic remainder in D moves the limiting simplex boundary by $O(\delta) = o(r^{-1})$. Integrating the density expansion and using

$$\int_{\{\sum_i a_i x_i \leq 1\}} \prod_i x_i^{\nu_i-1} dx = \frac{\prod_i \Gamma(\nu_i) a_i^{-\nu_i}}{\Gamma(\kappa + 1)}$$

gives (C.12). The factor multiplying z^κ is slowly varying and obeys (C.11), so Proposition C.5 applies. A positive weight $w(u) = w(0) + O(\|u\|_1)$ changes the rescaled integrand by a relative $O(\delta)$ and therefore only multiplies the leading constant. $\square$

For the AR(1) covariance of the four active Gaussian coordinates at $\rho = 1/4$, the corner exponents are

$$\kappa_j = \begin{cases} 34/15, & j = 1, 4, \\ 41/15, & j = 2, 3, \\ 14/5 + \lambda_j^2/(1 - \lambda_j^2), & j \notin \{1, 2, 3, 4\}, \end{cases} \quad \lambda_j = \rho^{j-4} \ (j \geq 5), \quad (\text{C.13})$$

so $34/15 \leq \kappa_j \leq 43/15$ uniformly; in B1 a proxy coordinate is independent of the active block and has $\kappa_j = 14/5$. All components of $\Omega \mathbf{1}$ are bounded away from zero. The reciprocal deficit has the required expansion: in A1 and A2 this follows by expanding the exponential index surface at the active corner, and in A3 the relevant first derivatives of its exponent are $0.8 + 0.20u$, $0.8 + 0.15u$ or 0.8 , hence uniformly at least 0.8 . In the A family the weight c is identically one; in B1 the bounded latent scale factor, integrated over the conditional Gaussian law of F , has a uniform expansion $w_{j,u}(z) = w_{j,u}(0) + O(\|z\|_1)$ with $0 < w_- \leq w_{j,u}(0) \leq w_+$, so the weighted clause of Lemma C.6 applies.

The remaining primitive assumptions hold as follows. The conditional generators are strictly increasing and atomless, which gives (P1). For A1–A3, $q(s, u) = s^{\gamma(u)} / [1 + \exp\{v(u) - s\}]$, and both the survival remainder and its log-level derivative decay faster than every exponential in $t = \log y$; for B1, $q(s, u, F) = s^{\gamma(u)} \exp\{A(F) - 1/(2s)\}$ with $A(F) = \{\Phi(F) - 1/2\}/2$, and inversion gives $\mathbb{P}(Y > y \mid U = x) = c_B(x) y^{-1/\gamma(x)} \{1 + O(y^{-2})\}$ with $c_B(x) = \mathbb{E}[e^{A(F)/\gamma} \mid U = x]$, the same order

holding for the log-level derivative. This is (P2). A common-residual coupling of the conditional Gaussian laws moves the other PIT coordinates by $O(|u - v|)$ on $I_{\varepsilon/2}$, which gives the $1/\gamma$ clause of (P4); the c and $\log(1 + r)$ clauses follow by differentiating the bounded posterior exponential moments, with $\zeta_n = 0$. The projected profiles are exactly

$$\xi_j(u) = \begin{cases} \frac{1}{2}e^{-u}, & \text{A1, B1, } j \leq 4, \\ \frac{1}{2}e^{-u/2}, & \text{A2, } j \leq 4, \\ \frac{1}{2}e^{-0.8u}, & \text{A3, } j \leq 4, \\ \frac{1}{2}, & j > 4. \end{cases}$$

Corollary C.7 (The simulation models satisfy the estimation conditions). *For A1–A3 and B1, conditions (E1), (E2) and (E3[†]) hold along every sequence*

$$\alpha_n = n^{-a}, \quad h_n = \frac{1}{2}n^{-b}, \quad p_n \leq n^C,$$

with fixed $0 < a < 1$ and $b, C > 0$, with deterministic bias orders

$$\omega_\xi(1/n) = O(n^{-1}), \quad b_n(3\alpha_n) = O(1/\log n), \quad \omega_n^\dagger(C_0 h_n, \tau_n, \alpha_n) = O(n^{-b} \log n).$$

If in addition $b < 1/2$ and $a + b < 1$, then $\log(p_n n) = o(nh_n^2)$ and $\log(p_n n) = o(n\alpha_n h_n)$, so the population gaps of A1–A3 and B1 being fixed and positive, Theorems 3.1 and 3.2 apply. This covers all nine exponent pairs used by the aggregated screen: intersecting the nine separation events, every active coordinate has rank at most s and every inactive coordinate rank at least $s + 1$ at every setting, so taking the minimum of the nine ranks preserves the same strict separation, and a finite union bound makes the intersection probability tend to one.

Remark C.8 (Finite-design honesty). Corollary C.7 is asymptotic. It does not certify that the $o(\Delta_{\min})$ biases are numerically small at the design used in Section 4: there $\log(1/\alpha) \simeq 2.28$, the lower endpoint entering (E2) is only $x_\alpha = \log\{(3\alpha)^{-1}\} \simeq 1.18$, and $h \log(pn)$ is much larger than the population gaps under the available uniform upper bounds. The finite-sample tables are empirical evidence about that regime; they are not a numerical verification of (E2), (E3[†]) or of the concentration bound.

References

- [1] Anass Aghbalou, François Portier, Anne Sabourin, and Chen Zhou. Tail inverse regression: Dimension reduction for prediction of extremes. *Bernoulli*, 30(1):503–533, 2024. doi: 10.3150/23-BEJ1606.
- [2] Julyan Arbel, Stéphane Girard, and Hadrien Lorenzo. Shrinkage for extreme partial least-squares, 2024. arXiv:2403.09503v2.
- [3] Stéphane Boucheron, Gábor Lugosi, and Pascal Massart. *Concentration Inequalities: A Nonasymptotic Theory of Independence*. Oxford University Press, Oxford, 2013. doi: 10.1093/acprof:oso/9780199535255.001.0001.
- [4] Meryem Bousebata, Geoffroy Enjolras, and Stéphane Girard. Extreme partial least-squares. *Journal of Multivariate Analysis*, 194:105101, 2023. doi: 10.1016/j.jmva.2022.105101.
- [5] Abdelaati Daouia, Laurent Gardes, and Stéphane Girard. On kernel smoothing for extremal quantile regression. *Bernoulli*, 19(5B):2557–2589, 2013. doi: 10.3150/12-BEJ466.

- [6] Laurens de Haan and Ana Ferreira. *Extreme Value Theory: An Introduction*. Springer, New York, 2006. doi: 10.1007/0-387-34471-3.
- [7] Jianqing Fan and Jinchi Lv. Sure independence screening for ultrahigh dimensional feature space. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 70(5): 849–911, 2008. doi: 10.1111/j.1467-9868.2008.00674.x.
- [8] Laurent Gardes. Tail dimension reduction for extreme quantile estimation. *Extremes*, 21(1): 57–95, 2018. doi: 10.1007/s10687-017-0300-x.
- [9] Laurent Gardes and Alexandre Podgorny. Dimension reduction for the estimation of the conditional tail index. *Scandinavian Journal of Statistics*, 52(3):1444–1476, 2025. doi: 10.1111/sjos.12792.
- [10] Laurent Gardes and Gilles Stupfler. Estimation of the conditional tail index using a smoothed local Hill estimator. *Extremes*, 17(1):45–75, 2014. doi: 10.1007/s10687-013-0174-5.
- [11] Xuming He, Lan Wang, and Hyokyoung Grace Hong. Quantile-adaptive model-free variable screening for high-dimensional heterogeneous data. *The Annals of Statistics*, 41(1):342–369, 2013. doi: 10.1214/13-AOS1087.
- [12] Bruce M. Hill. A simple general approach to inference about the tail of a distribution. *The Annals of Statistics*, 3(5):1163–1174, 1975. doi: 10.1214/aos/1176343247.
- [13] Runze Li, Wei Zhong, and Liping Zhu. Feature screening via distance correlation learning. *Journal of the American Statistical Association*, 107(499):1129–1139, 2012. doi: 10.1080/01621459.2012.695654.
- [14] Pascal Massart. The tight constant in the Dvoretzky–Kiefer–Wolfowitz inequality. *The Annals of Probability*, 18(3):1269–1283, 1990. doi: 10.1214/aop/1176990746.
- [15] Michael Redmond and Alok Baveja. A data-driven software tool for enabling cooperative information sharing among police departments. *European Journal of Operational Research*, 141(3):660–678, 2002. doi: 10.1016/S0377-2217(01)00264-8.
- [16] Alfréd Rényi. On the theory of order statistics. *Acta Mathematica Academiae Scientiarum Hungaricae*, 4:191–231, 1953. doi: 10.1007/BF02127580.
- [17] Yuya Sasaki, Jing Tao, and Yulong Wang. High-dimensional tail index regression, 2026. arXiv:2403.01318v3.
- [18] Hansheng Wang and Chih-Ling Tsai. Tail index regression. *Journal of the American Statistical Association*, 104(487):1233–1240, 2009. doi: 10.1198/jasa.2009.tm08458.
- [19] Takuma Yoshida and Yuta Umezu. Sure independence screening for covariate-dependent extreme value index estimation, 2026. arXiv:2410.15705v3.